
Half Title Page



Title Page

LOC Page

Vince: to Riggins
Geraint: also, to Riggins



Contents

Foreword	ix
Preface	xi
Contributors	xiii
SECTION I Getting Started	
CHAPTER 1 ■ Introduction	3
1.1 WHO IS THIS BOOK FOR?	3
1.2 WHAT DO WE MEAN BY APPLIED MATHEMATICS?	3
1.3 WHAT IS OPEN SOURCE SOFTWARE	4
1.4 HOW TO GET THE MOST OUT OF THIS BOOK	4
SECTION II Probabilistic Modelling	
CHAPTER 2 ■ Markov Chains	9
2.1 PROBLEM	9
2.2 THEORY	9
2.3 SOLVING WITH PYTHON	11
2.4 SOLVING WITH R	18
2.5 RESEARCH	25
CHAPTER 3 ■ Discrete Event Simulation	27
3.1 PROBLEM	27
3.2 THEORY	28
3.2.1 Event Scheduling Approach	29
3.2.2 Process Based Simulation	30
3.3 SOLVING WITH PYTHON	30

viii ■ Contents

3.4	SOLVING WITH R	37
3.5	RESEARCH	43

SECTION III Dynamical Systems

CHAPTER	4 ■ Modelling with Differential Equations	47
---------	---	----

4.1	PROBLEM	47
4.2	THEORY	47
4.3	SOLVING WITH PYTHON	48
4.4	SOLVING WITH R	53
4.5	RESEARCH	56

CHAPTER	5 ■ Systems dynamics	57
---------	----------------------	----

5.1	PROBLEM	57
5.2	THEORY	57
5.3	SOLVING WITH PYTHON	60
5.4	SOLVING WITH R	68
5.5	RESEARCH	75

SECTION IV Optimisation

CHAPTER	6 ■ Linear programming	79
---------	------------------------	----

6.1	PROBLEM	79
6.2	THEORY	80
6.3	SOLVING WITH PYTHON	83
6.4	SOLVING WITH R	87
6.5	RESEARCH	96

	Bibliography	97
--	--------------	----

Foreword

This is the foreword



Preface

This is the preface.



Contributors

Michaél Aftosmis

NASA Ames Research Center
Moffett Field, California

Pratul K. Agarwal

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Sadaf R. Alam

Oak Ridge National Laboratory
Oak Ridge, Tennessee

Gabrielle Allen

Louisiana State University
Baton Rouge, Louisiana

Martin Sandve Alnæs

Simula Research Laboratory and University
of Oslo, Norway
Norway

Steven F. Ashby

Lawrence Livermore National Laboratory
Livermore, California

David A. Bader

Georgia Institute of Technology
Atlanta, Georgia

Benjamin Bergen

Los Alamos National Laboratory
Los Alamos, New Mexico

Jonathan W. Berry

Sandia National Laboratories
Albuquerque, New Mexico

Martin Berzins

University of Utah

Salt Lake City, Utah

Abhinav Bhatele

University of Illinois
Urbana-Champaign, Illinois

Christian Bischof

RWTH Aachen University
Germany

Rupak Biswas

NASA Ames Research Center
Moffett Field, California

Eric Bohm

University of Illinois
Urbana-Champaign, Illinois

James Bordner

University of California, San Diego
San Diego, California

Geörge Bosilca

University of Tennessee
Knoxville, Tennessee

Grèg L. Bryan

Columbia University
New York, New York

Marian Bubak

AGH University of Science and Technology
Kraków, Poland

Andrew Canning

Lawrence Berkeley National Laboratory
Berkeley, California

xiv ■ Contributors

Jonathan Carter

Lawrence Berkeley National Laboratory
Berkeley, California

Zizhong Chen

Jacksonville State University
Jacksonville, Alabama

Joseph R. Crobak

Rutgers, The State University of New
Jersey

Piscataway, New Jersey

Roxana E. Diaconescu

Yahoo! Inc.
Burbank, California

Roxana E. Diaconescu

Yahoo! Inc.
Burbank, California

I

Getting Started



Introduction

THANK you for starting to read this book. This book aims to bring together two fascinating topics:

- Problems that can be solved using mathematics;
- Software that is free to use and change.

What we mean by both of those things will become clear through reading this chapter and the rest of the book.

1.1 WHO IS THIS BOOK FOR?

Anyone who is interested in using mathematics and computers to solve problems will hopefully find this book helpful.

If you are a student of a mathematical discipline, a graduate student of a subject like operational research, a hobbyist who enjoys solving the travelling salesman problem or even if you get paid to do this stuff: this book is for you. We will introduce you to the world of open source software that allows you to do all these things freely.

If you are a student learning to write code, a graduate student using databases for their research, an enthusiast who programmes applications to help coordinate the neighbourhood watch, or even if you get paid to write software: this book is for you. We will introduce you to a world of problems that can be solved using your skill sets.

It would be helpful for the reader of this book to:

- Have access to a computer and be able to connect to the internet (at least once) to be able to download the relevant software.
- Be prepared to read some mathematics. Technically you do not need to understand the specific mathematics to be able to use the tools in this book. The topics covered use some algebra, calculus and probability.

1.2 WHAT DO WE MEAN BY APPLIED MATHEMATICS?

We consider this book to be a book on applied mathematics. This is not however a universal term, for some applied mathematics is the study of mechanics and involves

modelling projectiles being fired out of canons. We will use the term a bit more freely here and mean any type of real world problem that can be tackled using mathematical tools. This is sometimes referred to as operational research, operations research, mathematical modelling or indeed just mathematics.

One of the authors, Vince, used mathematics to plan the sitting plan at his wedding. Using a particular area of mathematics call graph theory he was able to ensure that everyone sat next to someone they liked and/or knew.

The other author, Geraint, used mathematics to find the best team of Pokemon. Using an area of mathematics call linear programming which is based on linear algebra he was able to find the best makeup of pokemon.

Here, applied mathematics is the type of mathematics that helps us answer questions that the real world asks.

1.3 WHAT IS OPEN SOURCE SOFTWARE

Strictly speaking open source software is software with source code that anyone can read, modify and improve. In practice this means that you do not need to pay to use it which is often one of the first attractions. This financial aspect can also be one of the reasons that someone will not use a particular piece of software due to a confusion between cost and value: if something is free is it really going to be any good?

In practice open source software is used all of the world and powers some of the most important infrastructure around. For example, one should never use any cryptographic software that is not open source: if you cannot open up and read things than you should not trust it (this is indeed why most cryptographic systems used are open source).

Today, open source software is a lot more than a licensing agreement: it is a community of practice. Bugs are fixed faster, research is implemented immediately and knowledge is spread more widely thanks to open source software. Bugs are fixed faster because anyone can read and inspect the source code. Most open source software projects also have a clear mechanisms for communicating with the developers and even reviewing and accepting code contributions from the general public. Research is implemented immediately because when new algorithms are discovered they are often added directly to the software by the researchers who found them. This all contributes to the spread of knowledge: open source software is the modern should of giants that we all stand on.

Open source software is software that, like scientific knowledge is not restricted in its use.

1.4 HOW TO GET THE MOST OUT OF THIS BOOK

The book itself is open source. You can find the source files for this book online at github.com/drvinceknight/ampwoss. There will will also find a number of *Jupyter notebooks* and *R markdown files* that include code snippets that let you follow along.

We feel that you can choose to read the book from cover to cover, writing out

the code examples as you go; or it could also be used as a reference text when faced with particular problem and wanting to know where to start.

The book is made up of 10 chapters that are paired in two 4 parts. Each part corresponds to a particular area of mathematics, for example “Emergent Behaviour”. Two chapters are paired together for each chapter, usually these two chapters correspond to the same area of mathematics but from a slightly different scale that correspond to different ways of tackling the problem.

Every chapter has the following structure:

1. Introduction - a brief overview of a given problem type. Here we will describe the problem at hand in general terms.
2. An Example problem. This will provide a tangible example problem that offers the reader some intuition for the rest of the discussion.
3. Solving with Python. We will describe the mathematical tools available to us in a programming language called Python to solve the problem.
4. Solving with R. Here we will do the same with the R programming language.
5. Brief theoretic background with pointers to reference texts. Some readers might like to delve in to the mathematics of the problem a bit further, we will include those details here.
6. Examples of research using these methods. Finally, some readers might even be interested in finding out a bit more of what mathematicians are doing on these problems. Often this will include some descriptions of the problem considered but perhaps at a much larger scale than the one presented in the example.

For a given reader, not all sections of a chapter will be of interest. Perhaps a reader is only interested in R and finding out more about the research. Please do take from the book what you find useful.



II

Probabilistic Modelling



Markov Chains

MANY real world situations have some level of unpredictability through randomness: the flip of a coin, the number of orders of coffee in a shop, the winning numbers of the lottery. However, mathematics can in fact let us make predictions about what we expect to happen. One tool used to understand randomness is Markov chains, an area of mathematics sitting at the intersection of probability and linear algebra.

2.1 PROBLEM

Consider a barber shop. The shop owners have noticed that customers will not wait if there is no room in their waiting room and will choose to take their business elsewhere. The Barber shop would like to make an investment so as to avoid this situation. They know the following information:

- They currently have 2 barber chairs (and 2 barbers).
- They have waiting room for 4 people.
- They usually have 10 customers arrive per hour.
- Each Barber takes about 15 minutes to serve a customer so they can serve 4 customers an hour.

This is represented diagrammatically in Figure 2.1.

They are planning on reconfiguring space to either have 2 extra waiting chairs or another barber's chair and barber.

The mathematical tool used to model this situation is a Markov chain.

2.2 THEORY

A Markov chain is a model of a sequence of random events that is defined by a collection of **states** and rules that define how to move between these states.

For example, in the barber shop a single number is sufficient to describe the status of the shop. If that number is 1 this implies that 1 customer is currently having their

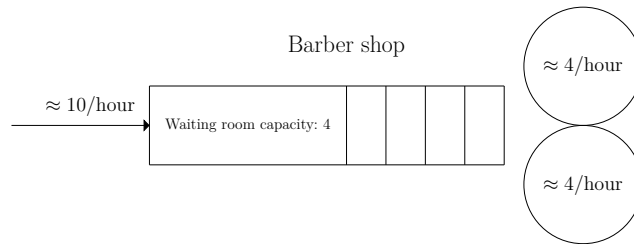


Figure 2.1 Diagrammatic representation of the barber shop as a queuing system.

hair cut. If that number is 5 this implies that 2 customers are being served and 3 are waiting. The entire state space is, in this case a finite set of integers from 0 to 6. If the system is full (all barbers and waiting room occupied) then we are in state 6 and if there is no one at the shop then we are in state 0. This is denoted mathematically as:

$$S = \{0, 1, 2, 3, 4, 5, 6\} \quad (2.1)$$

As customers arrive and leave the system goes between states as shown in Figure 2.2.




Figure 2.2 Diagrammatic representation of the state space

The rules that govern how to move between these states can be defined in two ways:

- Using probabilities of changing state (or not) in a well defined time period. This is called a discrete Markov chain.
- Using rates of change from one state to another. This is called a continuous time Markov chain.

For our barber shop we will consider it as a continuous Markov chain as shown in Figure 2.3

Note that a Markov chain assumes the rates follow an exponential distribution. One interesting property of this distribution is that it is considered memoryless which means that if a customer has been having their hair cut for 5 minutes this does not change the rate at which their service ends. This distribution is quite common in the real world and therefore a common assumption.

These states and rates can be represented mathematically using a transition matrix Q where Q_{ij} represents the rate of going from state i to state j . In this case we have:




Figure 2.3 Diagrammatic representation of the state space and the transition rates

$$Q = \begin{pmatrix} -10 & 10 & 0 & 0 & 0 & 0 & 0 \\ 4 & -14 & 10 & 0 & 0 & 0 & 0 \\ 0 & 8 & -18 & 10 & 0 & 0 & 0 \\ 0 & 0 & 8 & -18 & 10 & 0 & 0 \\ 0 & 0 & 0 & 8 & -18 & 10 & 0 \\ 0 & 0 & 0 & 0 & 8 & -18 & 10 \\ 0 & 0 & 0 & 0 & 0 & 8 & -8 \end{pmatrix} \quad (2.2)$$

You will see that Q_{ii} are negative and ensure the rows of Q sum to 0. This gives the total rate of change leaving state i .

We can use Q to understand the probability of being in a given state after t time units. This can be represented mathematically using a matrix P_t where $(P_t)_{ij}$ is the probability of being in state j after t time units having started in state i . We can use Q to calculate P_t using the matrix exponential:

$$P_t = e^{Qt} \quad (2.3)$$

What is also useful is understanding the long run behaviour of the system. This allows us to answer questions such as “what state are we most likely to be in on average?” or “what is the probability of being in the last state on average?”.

This long run probability distribution over the state can be represented using a vector π where π_i represents the probability of being in state i . This vector is in fact the solution to the following matrix equation:

$$\pi Q = 0 \quad (2.4)$$

In the upcoming sections we will demonstrate all of the above concepts.

2.3 SOLVING WITH PYTHON

The first step we will take is to write a function to obtain the transition rates between two given states:

Python input

```

1 def get_transition_rate(
2     in_state, out_state, waiting_room=4, num_barbers=2,
3 ):
4     """Return the transition rate for two given states.
5
6     Args:
7         in_state: an integer
8         out_state: an integer
9         waiting_room: an integer (default: 4)
10        num_barbers: an integer (default: 2)
11
12    Returns:
13        A real.
14    """
15    arrival_rate = 10
16    service_rate = 4
17
18    capacity = waiting_room + num_barbers
19    delta = out_state - in_state
20
21    if delta == 1 and in_state < capacity:
22        return arrival_rate
23
24    if delta == -1:
25        return min(in_state, num_barbers) * service_rate
26
27    return 0

```

Next, we write a function that creates an entire transition rate matrix Q for a given problem. We will use the `numpy` to handle all the linear algebra and the `itertools` library for some iterations:

Python input

```

28 import itertools
29 import numpy as np
30
31
32 def get_transition_rate_matrix(waiting_room=4, num_barbers=2):
33     """Return the transition matrix Q.
34
35     Args:
36         waiting_room: an integer (default: 4)
37         num_barbers: an integer (default: 2)
38
39     Returns:
40         A matrix.
41     """
42     capacity = waiting_room + num_barbers
43     state_pairs = itertools.product(
44         range(capacity + 1), repeat=2
45     )
46
47     flat_transition_rates = [
48         get_transition_rate(
49             in_state=in_state,
50             out_state=out_state,
51             waiting_room=waiting_room,
52             num_barbers=num_barbers,
53         )
54         for in_state, out_state in state_pairs
55     ]
56     transition_rates = np.reshape(
57         flat_transition_rates, (capacity + 1, capacity + 1)
58     )
59     np.fill_diagonal(
60         transition_rates, -transition_rates.sum(axis=1)
61     )
62
63     return transition_rates

```

Using this we can obtain the matrix Q for our default system:

Python input

```

64 Q = get_transition_rate_matrix()
65 print(Q)

```

which gives:

Python output

```

66 [[-10  10   0   0   0   0   0]
67  [  4 -14  10   0   0   0   0]
68  [  0   8 -18  10   0   0   0]
69  [  0   0   8 -18  10   0   0]
70  [  0   0   0   8 -18  10   0]
71  [  0   0   0   0   8 -18  10]
72  [  0   0   0   0   0   8 -8]]

```

We can take the matrix exponential as discussed above. To do this, we need to use the `scipy` library. To see what would happen after .5 time units we obtain:

Python input

```

73 import scipy.linalg
74
75 print(scipy.linalg.expm(Q * 0.5).round(5))

```

which gives:

Python output

```

76 [[0.10492 0.21254 0.20377 0.17142 0.13021 0.09564 0.0815 ]
77  [0.08501 0.18292 0.18666 0.1708  0.14377 0.1189  0.11194]
78  [0.06521 0.14933 0.16338 0.16478 0.15633 0.14751 0.15346]
79  [0.04388 0.10931 0.13183 0.15181 0.16777 0.18398 0.21142]
80  [0.02667 0.07361 0.10005 0.13422 0.17393 0.2189  0.27262]
81  [0.01567 0.0487  0.07552 0.11775 0.17512 0.24484 0.32239]
82  [0.01068 0.03668 0.06286 0.10824 0.17448 0.25791 0.34914]]

```

To see what would happen after 500 time units we obtain:

Python input

```
83 print(scipy.linalg.expm(Q * 500).round(5))
```

which gives:

Python output

```
84 [[0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
85  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
86  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
87  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
88  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
89  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]
90  [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]]
```

We see that no matter what state (row) the system is in, after 500 time units the probabilities are all the same. We could in fact stop our analysis here, however our choice of 500 time units was arbitrary and might not be the correct amount for all possible scenarios, as such we will continue to aim to solve the underlying equation 2.4 directly.

To do this we will solve the underlying system using a numerically efficient algorithm called least squares optimisation (available from the `numpy` library):

Python input

```

91 def get_steady_state_vector(Q):
92     """Return the steady state vector of any given continuous
93     time transition rate matrix.
94
95     Args:
96     Q: a transition rate matrix
97
98     Returns:
99     A vector
100     """
101     state_space_size, _ = Q.shape
102     A = np.vstack((Q.T, np.ones(state_space_size)))
103     b = np.append(np.zeros(state_space_size), 1)
104     x, _, _, _ = np.linalg.lstsq(A, b, rcond=None)
105     return x

```

So if we now see the steady state vector for our default system:

Python input

```

106 print(get_steady_state_vector(Q).round(5))

```

we get:

Python output

```

107 [0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176]

```

We can see that the shop is expected to be empty approximately 3.4% of the time and full 26.2% of the time.

The final function we will write is one that uses all of the above to just return the probability of the shop being full.

Python input

```

108 def get_probability_of_full_shop(
109     waiting_room=4, num_barbers=2
110 ):
111     """Return the probability of the barber shop being full.
112
113     Args:
114         waiting_room: an integer (default: 4)
115         num_barbers: an integer (default: 2)
116
117     Returns:
118         A real.
119     """
120     Q = get_transition_rate_matrix(
121         waiting_room=waiting_room, num_barbers=num_barbers,
122     )
123     pi = get_steady_state_vector(Q)
124     return pi[-1]

```

We can now confirm the previous probability calculated probability of the shop being full:

Python input

```

125 print(round(get_probability_of_full_shop(), 6))

```

which gives:

Python output

```

126 0.261756

```

If we were too have 2 extra space in the waiting room:

Python input

```

127 print(round(get_probability_of_full_shop(waiting_room=6), 6))

```

which gives:

Python output

```
0.23557
```

This is a slight improvement however, increasing the number of barbers has a substantial effect:

Python input

```
print(round(get_probability_of_full_shop(num_barbers=3), 6))
```

Python output

```
0.078636
```

2.4 SOLVING WITH R

The first step we will take is write a function to obtain the transition rates between two given states:

R input

```

131 #' Return the transition rate for two given states.
132 #'
133 #' @param in_state an integer
134 #' @param out_state an integer
135 #' @param waiting_room an integer (default: 4)
136 #' @param num_barbers an integer (default: 2)
137 #'
138 #' @return A real
139 get_transition_rate <- function(in_state,
140                                out_state,
141                                waiting_room = 4,
142                                num_barbers = 2){
143
144   arrival_rate <- 10
145   service_rate <- 4
146
147   capacity <- waiting_room + num_barbers
148   delta <- out_state - in_state
149
150   if (delta == 1) {
151     if (in_state < capacity) {
152       return(arrival_rate)
153     }
154   }
155
156   if (delta == -1) {
157     return(min(in_state, num_barbers) * service_rate)
158   }
159   return(0)
160 }

```

We will not actually use this function but a vectorized version of this:

R input

```

160 vectorized_get_transition_rate <- Vectorize(
161   get_transition_rate,
162   vectorize.args = c("in_state", "out_state")
163 )

```

This function can now take a vector of inputs for the `in_state` and `out_state` variables which will allow us to simplify the following code that creates the matrices:

R input

```

164  #' Return the transition rate matrix Q
165  #'
166  #' @param waiting_room an integer (default: 4)
167  #' @param num_barbers an integer (default: 2)
168  #'
169  #' @return A matrix
170  get_transition_rate_matrix <- function(waiting_room = 4,
171                                       num_barbers = 2){
172    max_state <- waiting_room + num_barbers
173
174    Q <- outer(0:max_state,
175              0:max_state,
176              vectorized_get_transition_rate,
177              waiting_room = waiting_room,
178              num_barbers = num_barbers
179            )
180    row_sums <- rowSums(Q)
181
182    diag(Q) <- -row_sums
183    Q
184  }

```

Using this we can obtain the matrix Q for our default system:

R input

```

185  Q <- get_transition_rate_matrix()
186  print(Q)

```

which gives:

R output

```

187      [,1] [,2] [,3] [,4] [,5] [,6] [,7]
188 [1,]  -10  10   0   0   0   0   0
189 [2,]   4 -14  10   0   0   0   0
190 [3,]   0  8 -18  10   0   0   0
191 [4,]   0  0  8 -18  10   0   0
192 [5,]   0  0  0  8 -18  10   0
193 [6,]   0  0  0  0  8 -18  10
194 [7,]   0  0  0  0  0  8 -8

```

One immediate thing we can do with this matrix is take the matrix exponential discussed above. To do this, we need to use an R library call `expm`.

To be able to make use of the nice `%>%` "pipe" operator we are also going to load the `magrittr` library. Now if we wanted to see what would happen after .5 time units we obtain:

R input

```

195 library(expm, warn.conflicts = FALSE, quietly = TRUE)
196 library(magrittr, warn.conflicts = FALSE, quietly = TRUE)
197
198 print( (Q * .5) %>% expm %>% round(5))

```

which gives:

R output

```

199      [,1] [,2] [,3] [,4] [,5] [,6] [,7]
200 [1,] 0.10492 0.21254 0.20377 0.17142 0.13021 0.09564 0.08150
201 [2,] 0.08501 0.18292 0.18666 0.17080 0.14377 0.11890 0.11194
202 [3,] 0.06521 0.14933 0.16338 0.16478 0.15633 0.14751 0.15346
203 [4,] 0.04388 0.10931 0.13183 0.15181 0.16777 0.18398 0.21142
204 [5,] 0.02667 0.07361 0.10005 0.13422 0.17393 0.21890 0.27262
205 [6,] 0.01567 0.04870 0.07552 0.11775 0.17512 0.24484 0.32239
206 [7,] 0.01068 0.03668 0.06286 0.10824 0.17448 0.25791 0.34914

```

After 500 time units we obtain:

R input

```
207 print( (Q * 500) %>% expm %>% round(5))
```

which gives:

R output

```
208      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
209 [1,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
210 [2,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
211 [3,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
212 [4,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
213 [5,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
214 [6,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
215 [7,] 0.03431 0.08577 0.10722 0.13402 0.16752 0.2094 0.26176
```

We see that no matter what state (row) we are in, after 500 time units the probabilities are all the same. We could in fact stop our analysis here, however our choice of 500 time units was arbitrary and might not be the correct amount for all possible scenarios, as such we will continue to aim to solve the underlying equation 2.4 directly.

To be able to do this, we will make use of the versatile **pracma** package which includes a number of numerical analysis functions for efficient computations.

R input

```

216 library(pracma, warn.conflicts = FALSE, quietly = TRUE)
217
218 #' Return the steady state vector of any given continuous time
219 #' transition rate matrix
220 #'
221 #' @param Q a transition rate matrix
222 #'
223 #' @return A vector
224 get_steady_state_vector <- function(Q){
225   state_space_size <- dim(Q)[1]
226   A <- rbind(t(Q), 1)
227   b <- c(integer(state_space_size), 1)
228   mldivide(A, b)
229 }

```

This is making use of `pracma`'s `mldivide` function which chooses the best numerical algorithm to find the solution to a given matrix equation $Ax = b$.

So if we now see the steady state vector for our default system:

R input

```

230 print(get_steady_state_vector(Q))

```

we get:

R output

```

231           [,1]
232 [1,] 0.03430888
233 [2,] 0.08577220
234 [3,] 0.10721525
235 [4,] 0.13401906
236 [5,] 0.16752383
237 [6,] 0.20940479
238 [7,] 0.26175598

```

We can see that the shop is expected to be empty approximately 3.4% of the time and full 26.2% of the time.

The final piece of this puzzle is to create a single function that uses all of the above to just return the probability of the shop being full.

R input

```

239 #' Return the probability of the barber shop being full
240 #'
241 #' @param waiting_room (default: 4)
242 #' @param num_barbers (default: 2)
243 #'
244 #' @return A real
245 get_probability_of_full_shop <- function(waiting_room = 4,
246                                         num_barbers = 2){
247   arrival_rate <- 10
248   service_rate <- 4
249   pi <- get_transition_rate_matrix(
250     waiting_room = waiting_room,
251     num_barbers = num_barbers
252   ) %>%
253     get_steady_state_vector()
254
255   capacity <- waiting_room + num_barbers
256   pi[capacity + 1]
257 }

```

Now we can run this code efficiently with both scenarios:

R input

```

258 print(get_probability_of_full_shop(waiting_room = 6))

```

which decreases the probability of a full shop to:

R output

```

259 [1] 0.2355699

```

but adding another barber and chair:

R input

260

```
print(get_probability_of_full_shop(num_barbers = 3))
```

gives:

R output

261

```
[1] 0.0786359
```

2.5 RESEARCH

TBA



Discrete Event Simulation

COMPLEX situations further compounded by randomness appear throughout our daily lives. For example, data flowing through a computer network, patients being treated at an emergency services, and daily commutes to work. Mathematics can be used to understand these complex situations so as to make predictions which in turn can be used to make improvements. One tool used to do this is to let a computer create a dynamic virtual representation of the scenario in question, the particular type we are going to cover here is called Discrete Event Simulation.

3.1 PROBLEM

Consider the following situation: a bicycle repair shop would like reconfigure their set-up in order to guarantee that all bicycles processed by the repair shop take a maximum of 30 minutes. Their current set-up is as follows:

- Bicycles arrive randomly at the shop at a rate of 15 per hour.
- They wait in line to be seen at an inspection counter, manned by one member of staff who can inspect one bicycle at a time. On average an inspection takes around 3 minutes.
- After inspection it is found that around 20% of bicycles do not need repair, and they are then ready for collection.
- After inspection it is found that around 80% of bicycles go on to be repaired. These then wait in line outside the repair workshop, which is manned by two members of staff who can each repair one bicycle at a time. On average a repair takes around 6 minutes.
- After repair the bicycles are ready for collection.

A diagram of the system is shown in Figure 3.1

We can also assume that there is infinite capacity at the bicycle repair shop for waiting bicycles. The shop will hire an extra member of staff in order to meet their target of a maximum time in the system of 30 minutes. They would like to know if they should work on the inspection counter or in the repair workshop?

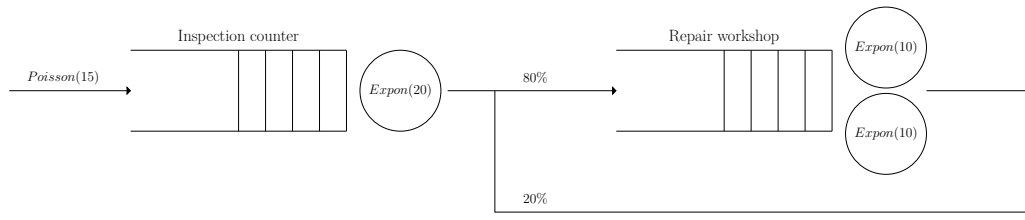


Figure 3.1 Diagrammatic representation of the bicycle repair shop as a queuing system.

3.2 THEORY

A number of the events that govern the behaviour of the bicycle shop above are probabilistic. For example the times that bicycles arrive at the shop, the duration of the inspection and repairs, and whether the bicycle would need to go on to be repaired or not. When a number of these probabilistic events are arranged in a complex system such as the bicycle shop, using analytical methods to manipulate these probabilities can become difficult. One method to deal with this is *simulation*.

Consider one probabilistic event, rolling a die. A die has six sides numbered 1 to 6, each side is equally likely to land. Therefore the probability of rolling a 1 is $\frac{1}{6}$, the probability of rolling a 2 is $\frac{1}{6}$, and so on. This means that that if we roll the die a large number of times, we would expect $\frac{1}{6}$ of those rolls to be a 1. This is called the *law of large numbers*.

Now imagine we have an event in which we do not know the analytical probability of it occurring. Consider rolling a weighted die, in this case a die in which the probability of obtaining one number is much greater than the others. How can we estimate the probability of obtaining a 5 on this die?

Rolling the weighted die once does not give us much information. However due to the law of large numbers, we can roll this die a number of times, and find the proportion of those rolls which gave a 5. The more times we roll the die, the closer this proportion approaches the underlying probability of obtaining a 5.

For a complex system such as the bicycle shop, we would like to estimate the proportion of bicycles that take longer than 30 minutes to be processed. As it is a complex system it is difficult to work this out analytically. So, just like the weighted die, we would like to observe this system a number of times and record the overall proportions of bicycles spending longer than 30 minutes in the shop, which will converge to the true underlying proportion. However unlike rolling a weighted die, it is costly to observe this shop over a number of days with identical conditions. In this case it is costly in terms of time, as the repair shop already exists. However some scenarios, for example the scenario where the repair shop hires an additional member of staff, do not yet exist, so observing this this would be costly in terms of money also. We can however build a virtual representation of this complex system on a computer, and observe a virtual day of work much more quickly and much less costly on the computer, similar to a video game.

In order to do this, the computer needs to be able to generate random outcomes of

each of the smaller events that make up the large complex system. Generating random events are essentially doing things to random numbers, that need to be generated.

Computers are deterministic, therefore true randomness is not always possible. They can however generate pseudorandom numbers: sequences of numbers that look like random numbers, but are entirely determined from the previous numbers in the sequence. Most programming languages have methods of doing this.

In order to simulate an event we can again manipulate the law of large numbers. Let $X \sim U(0, 1)$, a uniformly pseudorandom variable between 0 and 1. Let D be the outcome of a roll of an unbiased die. Then D can be defined as:

$$D = \begin{cases} 1 & \text{if } 0 \leq X < \frac{1}{6} \\ 2 & \text{if } \frac{1}{6} \leq X < \frac{2}{6} \\ 3 & \text{if } \frac{2}{6} \leq X < \frac{3}{6} \\ 4 & \text{if } \frac{3}{6} \leq X < \frac{4}{6} \\ 5 & \text{if } \frac{4}{6} \leq X < \frac{5}{6} \\ 6 & \text{if } \frac{5}{6} \leq X < 1 \end{cases} \quad (3.1)$$

The bicycle repair shop is a system made up of interactions of a number of other simpler random events. This can be thought of as many interactions of random variables, each generated using pseudorandom numbers.

In this case the fundamental random events that need to be generated are:

- the time each bicycle arrives to the repair shop,
- the time each bicycle spends at the inspection counter,
- whether each bicycle needs to go on the the repair workshop,
- the time each those bicycles spends at the repair shop.

As the simulation progresses these events will be generated, and will interact together as described in Section 6.1. The proportion of customers spending longer than 30 minutes in the shop can then be counted. This proportion itself is a random variable, and so just like the weighted die, running this simulation once does not give us much information. But we can run the simulation many times and take an average proportion, to smooth out any variability.

The process outlined above is a particular implementation of Monte Carlo simulation called *discrete event simulation*, which generates pseudorandom numbers and observes their interactions. In practice there are two main approaches to simulating complex probabilistic systems such as this one: the *event scheduling* approach, and *process based* simulation. It just so happens that the main implementations in Python and R use each of these approaches, so you will see both approaches used here.

3.2.1 Event Scheduling Approach

When using the event scheduling approach, we can think of the ‘virtual representation’ of the system as being the facilities that the bicycles use, shown in Figure 3.1.

Then we let entities (the bicycles) interact with these facilities. It is these facilities that determine how the entities behave.

In a simulation that uses an event scheduling approach, a key concept is that events occur that cause further events to occur in the future, either immediately or after a delay, such as after some time in service. In the bicycle shop examples of such events include a bicycle joining a queue, a bicycle beginning service, and a bicycle finishing service. At each event the event list is updated, and the clock then jumps forward to the next event in this updated list.

3.2.2 Process Based Simulation

When using process based simulation, we can think of the ‘virtual representation’ of the system as being the sequence of actions that each entity (the bicycles) must take, and these sequences of actions might contain delays as a number of entities seize and release a finite amount of resources. It is the sequence of actions that determine how the entities behave.

For the bicycle repair shop an example of one possible sequence of actions would be:

arrive → *seize inspection counter* → *delay* → *release inspection counter* → *seize repair shop* → *delay* → *release repair shop* → *leave*

The scheduled delays in this sequence of events correspond to the time spend being inspected and the time spend being repaired. Waiting in line for service at these facilities are not included in the sequence of events; that is implicit by the ‘seize’ and ‘release’ actions, as an entity will wait for a free resource before seizing one. Therefore in process based simulations, in addition to defining a sequence of events, resource types and their numbers also need to be defined.

3.3 SOLVING WITH PYTHON

In this book we will use the Ciw library in order to conduct discrete event simulation in Python. Ciw uses the event scheduling approach, which means we must define the system’s facilities, and then let customers loose to interact with them.

In this case there are two facilities to define: the inspection desk and the repair workshop. Let’s order these as so. For each of these we need to define:

- the distribution of times between consecutive bicycles arriving,
- the distribution of times the bicycles spend in service,
- the number of servers available,
- the probability of routing to each of the other facilities after service.

In this case we will assume that the time between consecutive arrivals follow a exponential distribution, and that the service times also follow exponential distributions. These are common assumptions for this sort of queueing system.

In Ciw, these are defined in a Network object, created using the `ciw.create_network`

function. The function below creates a Network object that defines the for a given set of parameters bicycle repair shop:

Python input

```

262 import ciw
263
264
265 def build_network_object(
266     num_inspectors=1, num_repairers=2,
267 ):
268     """Returns a Network object that defines the repair shop.
269
270     Args:
271         num_inspectors: a positive integer (default: 1)
272         num_repairers: a positive integer (default: 2)
273
274     Returns:
275         a Ciw network object
276     """
277     arrival_rate = 15
278     inspection_rate = 20
279     repair_rate = 10
280     prob_need_repair = 0.8
281     N = ciw.create_network(
282         arrival_distributions=[
283             ciw.dists.Exponential(arrival_rate),
284             ciw.dists.NoArrivals(),
285         ],
286         service_distributions=[
287             ciw.dists.Exponential(inspection_rate),
288             ciw.dists.Exponential(repair_rate),
289         ],
290         number_of_servers=[num_inspectors, num_repairers],
291         routing=[[0.0, prob_need_repair], [0.0, 0.0]],
292     )
293     return N

```

A Network object is used by Ciw to access system parameters. For example one piece of information it holds is the number of nodes of the system:

Python input

```

294 N = build_network_object()
295 print(N.number_of_nodes)

```

which gives:

Python output

```

296 2

```

Now we have defined the system, we need to use this to build the virtual representation of the system: in Ciw this is a Simulation object. Once this is built the simulation can be run, that is observe it for one virtual day. The following function does this:

Python input

```

297 def run_simulation(network, seed=0):
298     """Builds a simulation object and runs it for 8 time units.
299
300     Args:
301         network: a Ciw network object
302         seed: a float (default: 0)
303
304     Returns:
305         a Ciw simulation object after a run of the simulation
306     """
307     max_time = 8
308     ciw.seed(seed)
309     Q = ciw.Simulation(network)
310     Q.simulate_until_max_time(max_time)
311     return Q

```

Notice here a random seed is set. This is because there is some element of randomness when initialising this object, and much randomness in running the simulation, and in order to ensure reproducible results we force the pseudorandom number generator to produce the same sequence of pseudorandom numbers each time. Notice also that the simulation always begins with an empty system, so the first bicycle to arrive will never wait for service. Depending on the situation this may be an unwanted

feature, though not in this case as it is reasonable to assume that the bicycle shop will begin the day with no customers.

Now we wish to count the number of bicycles that have finished service, and to count the number of those whose entire journey through the system lasted longer than 0.5 hours. In order to do so, we'll use the `pandas` library for efficient manipulation of data frames.

Python input

```

312 import pandas as pd
313
314
315 def get_proportion(Q):
316     """Returns the proportion of bicycles spending over a given
317     limit at the repair shop.
318
319     Args:
320         Q: a Ciw simulation object after a run of the
321         simulation
322
323     Returns:
324         a real
325     """
326     limit = 0.5
327     inds = Q.nodes[-1].all_individuals
328     recs = pd.DataFrame(
329         dr for ind in inds for dr in ind.data_records
330     )
331     recs["total_time"] = (
332         recs["exit_date"] - recs["arrival_date"]
333     )
334     total_times = recs.groupby("id_number")["total_time"].sum()
335     return (total_times > limit).mean()

```

Altogether these functions can define the system, run one day of our system, and then find the proportion of bicycles spending over half an hour in the shop:

Python input

```
336 N = build_network_object()
337 Q = run_simulation(N)
338 p = get_proportion(Q)
339 print(round(p, 6))
```

This piece of code gives

Python output

```
340 0.261261
```

meaning 26.13% of all bicycles spent longer than half an hour at the repair shop.

However this particular day may have contained a number of extreme events. For a more accurate proportion this experiment should be repeated, and an average proportion taken. In order to do so, let's write a function that performs the above experiment over a number of trials, then finds an average proportion:

Python input

```

341 def get_average_proportion(num_inspectors=1, num_repairers=2):
342     """Returns the average proportion of bicycles spending over
343     a given limit at the repair shop.
344
345     Args:
346         num_inspectors: a positive integer (default: 1)
347         num_repairers: a positive integer (default: 2)
348
349     Returns:
350         a real
351     """
352     num_trials = 100
353     N = build_network_object(
354         num_inspectors=num_inspectors,
355         num_repairers=num_repairers,
356     )
357     proportions = []
358     for trial in range(num_trials):
359         Q = run_simulation(N, seed=trial)
360         proportion = get_proportion(Q=Q)
361         proportions.append(proportion)
362     return sum(proportions) / num_trials

```

This can be used to find the average proportion over 100 trials for the current system of one inspector and two repair people:

Python input

```

363 p = get_average_proportion(num_inspectors=1, num_repairers=2)
364 print(round(p, 6))

```

which gives:

Python output

```

365 0.159354

```

that is, on average 15.94% of bicycles will spend longer than 30 minutes at the repair shop.

Now consider the two possible future scenarios we wish to compare: hiring an extra member of staff to serve at the inspection desk, or hiring an extra member of staff at the repair workshop. Which scenario yields a smaller proportion of bicycles spending over 30 minutes at the shop? Let's investigate. First look at the situation where the additional member of staff works at the inspection desk:

Python input

```
366 p = get_average_proportion(num_inspectors=2, num_repairers=2)
367 print(round(p, 6))
```

which gives:

Python output

```
368 0.038477
```

that is 3.85% of bicycles.

Now look at the situation where the additional member of staff works at the repair workshop:

Python input

```
369 p = get_average_proportion(num_inspectors=1, num_repairers=3)
370 print(round(p, 6))
```

which gives:

Python output

```
371 0.103591
```

that is 10.36% of bicycles.

Therefore an additional member of staff at the inspection desk would be more beneficial than an additional member of staff at the repair workshop.

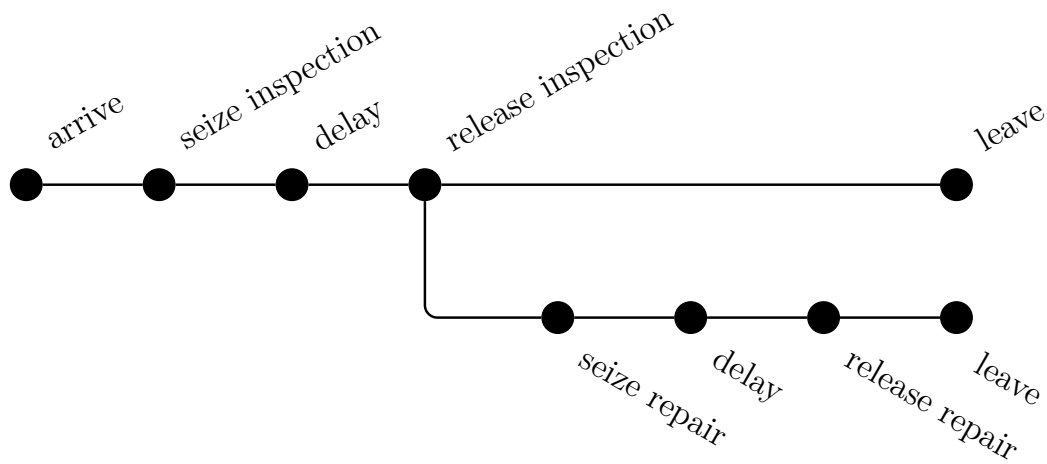


Figure 3.2 Diagrammatic representation of the forked trajectories a bicycle can take

3.4 SOLVING WITH R

In this book we will use the Simmer package in order to conduct discrete event simulation in R. Simmer uses the process based approach, which means we must define the each bicycle's sequence of actions, and then generate a number of bicycles with these sequences.

In Simmer these sequences of actions are made up of trajectories. The diagram in Figure 3.2 shows the branched trajectories than a bicycle would take at the repair shop:

The function below defines a simmer object that describes these trajectories:

R input

```

372 library(simmer)
373
374 #' Returns a simmer trajectory object outlining the bicycles
375 #' path through the repair shop
376 #'
377 #' @return A simmer trajectory object
378 define_bicycle_trajectories <- function() {
379   inspection_rate <- 20
380   repair_rate <- 10
381   prob_need_repair <- 0.8
382   bicycle <-
383     trajectory("Inspection") %>%
384     seize("Inspector") %>%
385     timeout(function() {
386       rexp(1, inspection_rate)
387     }) %>%
388     release("Inspector") %>%
389     branch(
390       function() (runif(1) < prob_need_repair),
391       continue = c(F),
392       trajectory("Repair") %>%
393         seize("Repairer") %>%
394         timeout(function() {
395           rexp(1, repair_rate)
396         }) %>%
397         release("Repairer"),
398       trajectory("Out")
399     )
400   return(bicycle)
401 }

```

These trajectories are not very useful alone, we are yet to define the resources used, or a way to generate bicycles with these trajectories. This is done in the function below, which begins by defining a `repair_shop` with one resource labelled “Inspector”, and two resources labelled “Repairer”. Once this is built the simulation can be run, that is observe it for one virtual day. The following function does all this:

R input

```

402 #' Runs one trial of the simulation.
403 #'
404 #' @param bicycle a simmer trajectory object
405 #' @param num_inspectors positive integer (default: 1)
406 #' @param num_repairers positive integer (default: 2)
407 #' @param seed a float (default: 0)
408 #'
409 #' @return A simmer simulation object after one run of
410 #;         the simulation
411 run_simulation <- function(bicycle,
412                           num_inspectors = 1,
413                           num_repairers = 2,
414                           seed = 0) {
415   arrival_rate <- 15
416   max_time <- 8
417   repair_shop <-
418     simmer("Repair Shop") %>%
419     add_resource("Inspector", num_inspectors) %>%
420     add_resource("Repairer", num_repairers) %>%
421     add_generator("Bicycle", bicycle, function() {
422       rexp(1, arrival_rate)
423     })
424
425   set.seed(seed)
426   repair_shop %>% run(until = 8)
427   return(repair_shop)
428 }

```

Notice here a random seed is set. This is because there are elements of randomness when running the simulation, and in order to ensure reproducible results we force the pseudorandom number generator to produce the same sequence of pseudorandom numbers each time. Notice also that the simulation always begins with an empty system, so the first bicycle to arrive will never wait for service. Depending on the situation this may be an unwanted feature, though not in this case as it is reasonable to assume that the bicycle shop will begin the day with no customers.

Now we wish to count the number of bicycles that have finished service, and to count the number of those whose entire journey through the system lasted longer than 0.5 hours. Using Simmer's `get_mon_arrivals()` function we can get a data frame of records to manipulate.

R input

```

429  #' Returns the proportion of bicycles spending over 30
430  #' minutes in the repair shop
431  #'
432  #' @param repair_shop a simmer simulation object
433  #'
434  #' @return a float between 0 and 1
435  get_proportion <- function(repair_shop) {
436    limit <- 0.5
437    recs <- repair_shop %>% get_mon_arrivals()
438    total_times <- recs$end_time - recs$start_time
439    return(mean(total_times > 0.5))
440  }

```

Altogether these functions can define the system, run one day of our system, and then find the proportion of bicycles spending over half an hour in the shop:

R input

```

441  bicycle <- define_bicycle_trajectories()
442  repair_shop <- run_simulation(bicycle = bicycle)
443  print(get_proportion(repair_shop = repair_shop))

```

This piece of code gives

R output

```

444  [1] 0.04032258

```

meaning 4.03% of all bicycles spent longer than half an hour at the repair shop.

However this particular day may have contained a number of extreme events. For a more accurate proportion this experiment should be repeated, and an average proportion taken. In order to do so, let's write a function that performs the above experiment over a number of trials, then finds an average proportion:

R input

```

445 #' Returns the average proportion of bicycles spending over
446 #' a given limit at the repair shop.
447 #'
448 #' @param num_inspectors positive integer (default: 1)
449 #' @param num_repairers positive integer (default: 2)
450
451 #' @return a float between 0 and 1
452 get_average_proportion <- function(num_inspectors = 1,
453                                   num_repairers = 2) {
454   num_trials <- 100
455   bicycle <- define_bicycle_trajectories()
456   proportions <- c()
457   for (trial in 1:num_trials) {
458     repair_shop <- run_simulation(
459       bicycle = bicycle,
460       num_inspectors = num_inspectors,
461       num_repairers = num_repairers,
462       seed = trial
463     )
464     proportion <- get_proportion(
465       repair_shop = repair_shop
466     )
467     proportions[trial] <- proportion
468   }
469   return(mean(proportions))
470 }

```

This can be used to find the average proportion over 100 trials:

R input

```

471 print(
472   get_average_proportion(
473     num_inspectors = 1,
474     num_repairers = 2)
475 )

```

which gives:

R output

476 [1] 0.1551579

that is, on average 15.52% of bicycles will spend longer than 30 minutes at the repair shop.

Now consider the two possible future scenarios we wish to compare: hiring an extra member of staff to serve at the inspection desk, or hiring an extra member of staff at the repair workshop. Which scenario yields a smaller proportion of bicycles spending over 30 minutes at the shop? Let's investigate. First look at the situation where the additional member of staff works at the inspection desk:

R input

```
477 print(
478   get_average_proportion(
479     num_inspectors = 2,
480     num_repairers = 2)
481 )
```

which gives:

R output

482 [1] 0.04115338

that is 4.12% of bicycles.

Now look at the situation where the additional member of staff works at the repair workshop:

R input

```
483 print(
484   get_average_proportion(
485     num_inspectors = 1,
486     num_repairers = 3)
487 )
```

which gives:

R output

488

```
[1] 0.1000899
```

that is 10.01% of bicycles.

Therefore an additional member of staff at the inspection desk would be more beneficial than an additional member of staff at the repair workshop.

3.5 RESEARCH

TBA



III

Dynamical Systems



Modelling with Differential Equations

SYSTEMS often change in a way that depends on their current state. For example, the speed at which a cup of coffee cools down depends on its current temperature. These types of systems are called dynamical systems and are modelled mathematically using differential equations. In this chapter we will consider a direct solution approach using symbolic mathematics.

4.1 PROBLEM

Consider the following situation: the entire population of a small rural town has caught a cold. All 100 individuals will recover at an average rate of 2 per day. The town leadership have noticed that being ill costs approximately £10 per day, this is due to general lack of productivity, poorer mood and other intangible aspects. They need to decide whether or not to order cold medicine which would **double** the recover rate. The cost of the cold medicine is a one off cost of £5 per person.

4.2 THEORY

In the case of this town, the overall rate at which people get better is dependent on the number of people in how are ill. This can be represented mathematically using a differential equation which is a way of relating the rate of change of a system to the state of the system itself.

In general if we are interested in some variable x over time t the differential function equation will be of the form:

$$\frac{dx}{dt} = f(x) \quad (4.1)$$

For some function f . In our case, if we denote the number of infected individuals as I where we implicitly mean that I is a function of time: $I = I(t)$ and the rate at which individuals recover by α then the differential equation that describes the above situation is:

$$\frac{dI}{dt} = -\alpha I \quad (4.2)$$

Finding a solution to this differential equation means finding an expression for I that when differentiated gives $-\alpha I$.

In this particular case, one such function is:

$$I(t) = e^{-\alpha t} \quad (4.3)$$

However, $I(0) = 1$ whereas for our problem we know that at time $t = 0$ there are 100 infected individuals. Indeed a differential equation defines a family of solutions and we need to know some sort of initial (also referred to as boundary) condition to have the exact solution. Which in this case would be:

$$I(t) = 100e^{-\alpha t} \quad (4.4)$$

To evaluate the cost we then need to know the sum of the values of that function over time. Integration gives us exactly this, so the cost would be:

$$K \int_0^{\infty} I(t) dt \quad (4.5)$$

where K is the cost per person per unit time.

In the upcoming sections we will confirm and use code to carry out the above efficiently so as to answer the original question.

4.3 SOLVING WITH PYTHON

The first step we will take is to write a function to obtain the differential equation. Note that here we will be using the Python library `sympy` which allows us to carry out symbolic calculations.

Python input

```

489 import sympy as sym
490
491 t = sym.Symbol("t")
492 alpha = sym.Symbol("alpha")
493 I_0 = sym.Symbol("I_0")
494 I = sym.Function("I")
495
496
497 def get_equation(alpha=alpha):
498     """Return the differential equation.
499
500     Args:
501         alpha: a float (default: symbolic alpha)
502
503     Returns:
504         A symbolic equation
505     """
506     return sym.Eq(sym.Derivative(I(t), t), -alpha * I(t))

```

Using this we can get the equation that defines the population change over time:

Python input

```

507 eq = get_equation()
508 print(eq)

```

which gives:

Python output

```

509 Eq(Derivative(I(t), t), -alpha*I(t))

```

Note that if you are using Jupyter then your output will actually be a well rendered mathematical equation:

$$\frac{d}{dt}I(t) = -\alpha I(t)$$

Note that we can pass a value to α if we want to:

Python input

```

510 eq = get_equation(alpha=1)
511 print(eq)

```

Python output

```

512 Eq(Derivative(I(t), t), -I(t))

```

We will now write a function to obtain the solution to this differential

Python input

```

513 def get_solution(I_0=I_0, alpha=alpha):
514     """Return the solution to the differential equation.
515
516     Args:
517         I_0: a float (default: symbolic I_0)
518         alpha: a float (default: symbolic alpha)
519
520     Returns:
521         A symbolic equation
522     """
523     eq = get_equation(alpha=alpha)
524     return sym.dsolve(eq, I(t), ics={I(0): I_0})

```

We can verify the solution discussed previously:

Python input

```

525 sol = get_solution()
526 print(sol)

```

which gives:

Python output

```
527 Eq(I(t), I_0*exp(-alpha*t))
```

$$I(t) = I_0 e^{-\alpha t}$$

We can use sympy itself to verify the result, by taking the derivative of the right hand side of our solution.

Python input

```
528 print(sym.diff(sol.rhs, t) == -alpha * sol.rhs)
```

which gives:

Python output

```
529 True
```

All of the above has given us the general solution in terms of $I(0) = I_0$ and α however we have written the code in such a way as we can pass the actual parameters:

Python input

```
530 sol = get_solution(alpha=2, I_0=100)
531 print(sol)
```

which gives:

Python output

```
532 Eq(I(t), 100*exp(-2*t))
```

Now, to calculate the cost we will write a function to integrate our result:

Python input

```

533 def get_cost(
534     I_0=I_0, alpha=alpha, cost_per_person=10, cost_of_cure=0,
535 ):
536     """Return the cost.
537
538     Args:
539         I_0: a float (default: symbolic I_0)
540         alpha: a float (default: symbolic alpha)
541         cost_per_person: a float (default: 10)
542         cost_of_cure: a float (default: 0)
543
544     Returns:
545         A symbolic expression
546     """
547     I_sol = get_solution(I_0=I_0, alpha=alpha)
548     return (
549         sym.integrate(I_sol.rhs, (t, 0, sym.oo))
550         * cost_per_person
551         + cost_of_cure * I_0
552     )

```

We can now obtain the cost without purchasing the cure:

Python input

```

553 I_0 = 100
554 alpha = 2
555 cost_without_cure = get_cost(I_0=I_0, alpha=alpha)
556 print(cost_without_cure)

```

which gives:

Python output

```

557 500

```

The cost with cure can use the above with a modified α and a non zero cost of the cure itself:

Python input

```

558 cost_of_cure = 5
559 cost_with_cure = get_cost(
560     I_0=I_0, alpha=2 * alpha, cost_of_cure=cost_of_cure
561 )
562 print(cost_with_cure)

```

which gives:

Python output

```

563 750

```

So given the current parameters it is not worth purchasing the cure.

4.4 SOLVING WITH R

R does not have a bespoke symbolic mathematics library however it has an interface to use Python's `sympy` called `rSymPy`. In general we would recommend to use Python directly when needing to carry out symbolic mathematical calculations however this is not always possible. Furthermore, the development of `rSymPy` seems to no longer be active and due to the dependency on Java could be problematic to use. One of other the difficulties of using `rSymPy` is that it works by passing string back and forth to Python's `sympy`. As such, we will need to have more of a manual approach and so we will not be able to write the modular code we have written throughout the book.

First let us setup and solve the differential equation:

R input

```

564 library("rSymPy", warn.conflicts = FALSE, quietly = TRUE)
565
566 t <- Var("t")
567 alpha <- Var("alpha")
568 I <- Var("I")
569 sympy("dsolve(Derivative(I(t), t) - alpha * I(t), I(t))")

```

this will give us:

R output

570

```
[1] "C1*exp(alpha*t)"
```

We can then solve the equation directly to find the initial condition:

R input

571

```
sympy("solve(C1*exp(alpha*0) - I_0, C1)")
```

This gives:

R output

572

```
[1] "[I_0]"
```

Note here that this is in fact using a much older version of Sympy so we cannot take advantage of some of the newer commands.

So we now know the form of our particular solution, thus we can now write a function that will give us the cost:

R input

```

573 #' Return the cost
574 #'
575 #' @param I_0 a float (default: symbolic I_0)
576 #' @param alpha a float (default: symbolic alpha)
577 #' @param cost_per_person a float (default: 10)
578 #' @param cost_of_cure a float (default: 0)
579 #'
580 #' @return A symbolic expression
581 get_cost <- function(I_0 = Var("I_0"),
582                      alpha = Var("alpha"),
583                      cost_per_person = 10,
584                      cost_of_cure = 0) {
585   t <- Var("t")
586   solution_cmd <- sprintf(
587     "I_sol = %s * exp(-%s * t)",
588     I_0,
589     alpha
590   )
591   sympy(solution_cmd)
592   cost_cmd <- sprintf(
593     "integrate(I_sol, (t, 0, oo)) * %s + %s * %s",
594     cost_per_person, cost_of_cure,
595     I_0
596   )
597   return(sympy(cost_cmd))
598 }

```

Using this:

R input

```

599 I_0 <- 100
600 alpha <- 2
601 cost_without_cure <- get_cost(I_0 = I_0, alpha = alpha)
602 print(cost_without_cure)

```

which gives:

R output

```
603 500
```

The cost with cure can use the above with a modified α and a non zero cost of the cure itself:

R input

```
604 cost_of_cure <- 5
605 cost_with_cure <- get_cost(
606     I_0 = I_0, alpha = 2 * alpha, cost_of_cure = cost_of_cure
607 )
608 print(cost_with_cure)
```

which gives:

R output

```
609 750
```

So given the current parameters it is not worth purchasing the cure.

4.5 RESEARCH

TBA

Systems Dynamics

IN many situations systems are dynamical, in that the state or population of a number of entities or classes change according to the current state or population of the system. For example population dynamics, chemical reactions, and systems of macroeconomics. It is often useful to be able to predict how these systems will behave over time, though the rules that govern these changes may be complex, and are not necessarily solvable analytically. In these cases numerical methods and visualisation may be used, which is the focus of this chapter.

5.1 PROBLEM

Consider the following scenario, where a population of 3000 people are susceptible to infection by some disease. This population can be described by the following parameters:

- They have a birth rate b of 0.01 per day;
- They have a death rate d of 0.01 per day;
- For every infectious individual, the infection rate α is 0.3 per day;
- Infectious people recover naturally (and thus gain an immunity from the disease), at a recovery rate r of 0.02 per day;
- For each day an individual is infected, they must take medication which costs a public healthcare system £10 per day.

A vaccine is produced, that allows new born individuals to gain an immunity. This vaccine costs the public health care system a one-off cost of £220 per vaccine. The healthcare providers would like to know if achieving a vaccination rate v of 85% would be beneficial financially.

5.2 THEORY

The above scenario is called a compartmental model of disease, and can be shown in the stock and flow diagram in Figure 5.1.

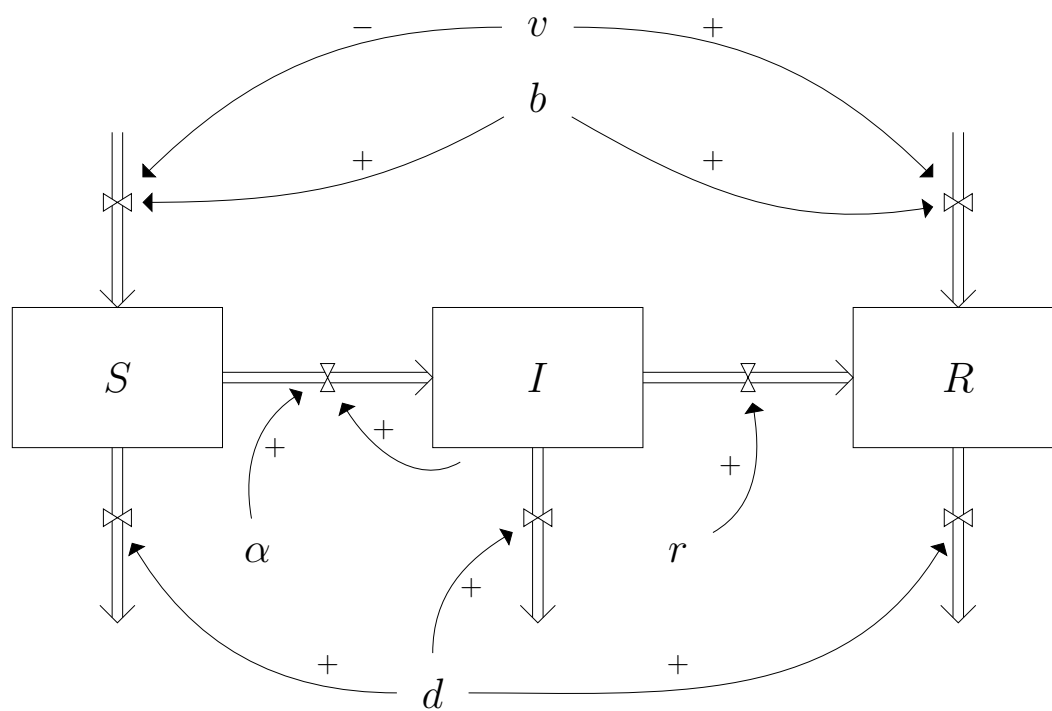


Figure 5.1 Diagrammatic representation of the epidemiology model

The system has three ‘stocks’ of different types of individuals, those susceptible to disease (S), those infected with the disease (I), and those who have recovered from the disease and so have gained immunity (R). The levels on these stocks change according to the flows in, out, and between them, controlled by ‘taps’. The amount of flow the taps let through are influenced in a multiplicative way (either negatively or positively), by other factors, such as external parameters (e.g. birth rate, infection rate) and the stock levels.

In this system the following taps exist, influenced by the following parameters:

- $external \rightarrow S$: Influenced positively by the birth rate, and negatively by the vaccine rate.
- $S \rightarrow I$: Influenced positively by the infection rate, and the number of infected individuals.
- $S \rightarrow external$: Influenced positively by the death rate.
- $I \rightarrow R$: Influenced positively by the recovery rate.
- $I \rightarrow external$: Influenced positively by the death rate.
- $R \rightarrow external$: Influenced positively by the birth rate and the vaccine rate.
- $external \rightarrow R$: Influenced positively by the death rate.

Mathematically the change in stock levels are written as the derivatives, for example the change in the number of susceptible individuals over time is denoted by $\frac{dS}{dt}$. This is equal to the sum of the taps in or out of that stock. Thus the system is described by the following system of differential equations:

$$\frac{dS}{dt} = -\frac{\alpha SI}{N} + (1-v)bN - dS \quad (5.1)$$

$$\frac{dI}{dt} = \frac{\alpha SI}{N} - (r+d)I \quad (5.2)$$

$$\frac{dR}{dt} = rI - dR + vbN \quad (5.3)$$

Where $N = S + I + R$ is the total number of individuals in the system.

We would like to understand the behaviour of the functions S , I and R under these rules, that is we would like to solve this system of differential equations. This system contains some non-linear terms, implying that this may be difficult to solve analytically, so we will use a numerical method instead.

There are a number of numerical methods, and the solvers we will use in Python and R cleverly choose the most appropriate for the problem at hand. In general methods for this kind of problems use the principle that the derivative denotes the rate of instantaneous change. Thus for a differential equation $\frac{dy}{dt} = f(t, y)$, consider the function y as a discrete sequence of points $\{y_0, y_1, y_2, y_3, \dots\}$ on $\{t_0, t_0 + h, t_0 + 2h, t_0 + 3h, \dots\}$ then

$$y_{n+1} = h \times f(t_0 + nh, y_n). \quad (5.4)$$

This sequence approaches the true solution y as $h \rightarrow 0$. Thus numerical methods, including the Runge-Kutta methods and the Euler method, step through this sequence $\{y_n\}$, choosing appropriate values of h and employing other methods of error reduction.

5.3 SOLVING WITH PYTHON

In this book we will use the `odeint` method of the SciPy library to numerically solve the above epidemiology models.

We first define the system of differential equations described in Equations 5.1, 5.2 and 5.3. This is a regular Python function, where the first two arguments are the system state and the current time respectively.

Python input

```

610 def derivatives(y, t, vaccine_rate, birth_rate=0.01):
611     """Defines the system of differential equations that
612     describe the epidemiology model.
613
614     Args:
615         y: a tuple of three integers
616         t: a positive float
617         vaccine_rate: a positive float <= 1
618         birth_rate: a positive float <= 1
619
620     Returns:
621         A tuple containing dS, dI, and dR
622     """
623     infection_rate = 0.3
624     recovery_rate = 0.02
625     death_rate = 0.01
626     S, I, R = y
627     N = S + I + R
628     dSdt = (
629         -((infection_rate * S * I) / N)
630         + ((1 - vaccine_rate) * birth_rate * N)
631         - (death_rate * S)
632     )
633     dIdt = (
634         ((infection_rate * S * I) / N)
635         - (recovery_rate * I)
636         - (death_rate * I)
637     )
638     dRdt = (
639         (recovery_rate * I)
640         - (death_rate * R)
641         + (vaccine_rate * birth_rate * N)
642     )
643     return dSdt, dIdt, dRdt

```

Using this function returns the instantaneous rate of change for each of the three stocks, S , I and R . If we begin at time 0.0, with 4 susceptible individuals, 1 infected individual, 0 recovered individuals, and a vaccine rate of 50%, then:

Python input

```
644 print(derivatives(y=(4, 1, 0), t=0.0, vaccine_rate=0.5))
```

Python output

```
645 (-0.255, 0.21, 0.045)
```

we would expect the number of susceptible individuals to reduce by around 0.255 per time unit, the number of infected individuals to increase by 0.21 per time unit, and the number of recovered individuals to increase by 0.045 per time unit. Now of course, after a tiny fraction of a time unit the stock levels will change, and thus the rates of change will change. So we will require something more sophisticated in order to determine the true behaviour of the system.

The following function observes the system's behaviour over some time period, using SciPy's `odeint` to numerically solve the system of differential equations:

Python input

```

646 from scipy.integrate import odeint
647
648
649 def integrate_ode(
650     derivative_function,
651     t,
652     y0=(2999, 1, 0),
653     vaccine_rate=0.85,
654     birth_rate=0.01,
655 ):
656     """Numerically solve the system of differential equations.
657
658     Args:
659         derivative_function: a function returning a tuple
660             of three floats
661         t: an array of increasing positive floats
662         y0: a tuple of three integers (default: (2999, 1, 0))
663         vaccine_rate: a positive float <= 1 (default: 0.85)
664         birth_rate: a positive float <= 1 (default: 0.01)
665
666     Returns:
667         A tuple of three arrays
668     """
669     results = odeint(
670         derivative_function,
671         y0,
672         t,
673         args=(vaccine_rate, birth_rate),
674     )
675     S, I, R = results.T
676     return S, I, R

```

Now we can use this function to investigate the difference in behaviour between a vaccination rate of 0% and a vaccination rate of 85%. Let's observe the system for two years, that is 730 days, in time steps of 0.01 days.

Begin with a vaccine rate of 0%:

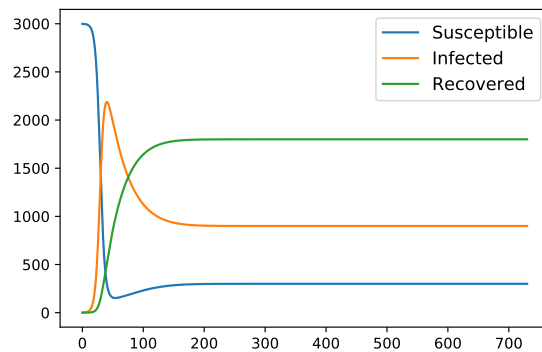


Figure 5.2 Output of code line 737-742

Python input

```

677 import numpy as np
678 from scipy.integrate import odeint
679
680 t = np.arange(0, 730.01, 0.01)
681 S, I, R = integrate_ode(derivatives, t, vaccine_rate=0.0)

```

Now S , I and R are arrays of values of the stock levels of S , I and R over the time steps t . Using `matplotlib` we can plot these to visualise their behaviour. The following code gives the plot shown in Figure 5.2.

Python input

```

682 import matplotlib.pyplot as plt
683
684 fig, ax = plt.subplots(1)
685 ax.plot(t, S, label='Susceptible')
686 ax.plot(t, I, label='Infected')
687 ax.plot(t, R, label='Recovered')
688 ax.legend(fontsize=12)
689 fig.savefig("plot_no_vaccine_python.pdf")

```

We observe that the number of infected individuals increases quickly, and in fact the rate of change increases as more individuals are infected. However this growth

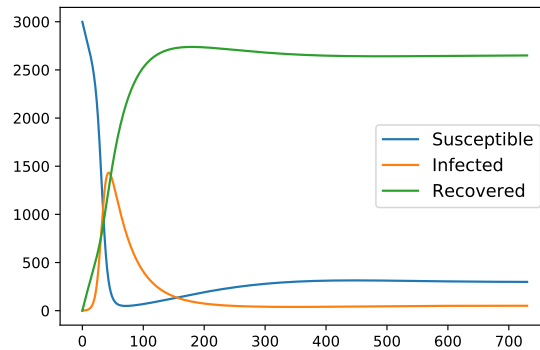


Figure 5.3 Output of code line 745-750

slows down as there are fewer susceptible individuals to infect. Due to the equal birth and death rates the overall population size remains constant; but we also see after some time period (around 300 time units) the levels of susceptible, infected, and recovered individuals becomes seemingly steady, and the disease becomes endemic. We can estimate once this steadiness occurs, around 10% of the population remain susceptible to the disease, 30% are infected, and 60% are recovered and immune.

Now with a vaccine rate of 85%:

Python input

```
690 t = np.arange(0, 730.01, 0.01)
691 S, I, R = integrate_ode(derivatives, t, vaccine_rate=0.85)
```

And again we can plot these to visualise their behaviour. The following code gives the plot shown in Figure 5.3.

Python input

```
692 fig, ax = plt.subplots(1)
693 ax.plot(t, S, label='Susceptible')
694 ax.plot(t, I, label='Infected')
695 ax.plot(t, R, label='Recovered')
696 ax.legend(fontsize=12)
697 fig.savefig("plot_with_vaccine_python.pdf")
```

With vaccination the disease remains endemic, however now we estimate that

once, steadiness occurs, around 10% of the population remain susceptible to the disease, 1.7% are infected, and 88.3% are immune or recovered and immune.

We've seen that vaccination lowers the percentage of the population living with the infection, which will lower the public healthcare system's medication costs. Let's now investigate if this saving is comparable to the cost of providing the vaccination to the newborns.

The following function calculates the total cost to the public healthcare system, that is the sum of the medication costs for those living with the infection and the vaccination costs:

Python input

```

698 def daily_cost(
699     derivative_function=derivatives, vaccine_rate=0.85
700 ):
701     """Calculates the daily cost to the public health system
702     after 2 years.
703
704     Args:
705         derivative_function: a function returning a tuple
706                             of three floats
707         vaccine_rate: a positive float <= 1 (default: 0.85)
708
709     Returns:
710         the daily cost
711     """
712     max_time = 730
713     time_step = 0.01
714     birth_rate = 0.01
715     vaccine_cost = 220
716     medication_cost = 10
717     t = np.arange(0, max_time + time_step, time_step)
718     S, I, R = integrate_ode(
719         derivatives,
720         t,
721         vaccine_rate=vaccine_rate,
722         birth_rate=birth_rate,
723     )
724     N = S[-1] + I[-1] + R[-1]
725     daily_vaccine_cost = (
726         N * birth_rate * vaccine_rate * vaccine_cost
727     ) / time_step
728     daily_meds_cost = (I[-1] * medication_cost) / time_step
729     return daily_vaccine_cost + daily_meds_cost

```

Now let's compare the total daily cost with and without vaccination. Without vaccinations:

Python input

```

730 cost = daily_cost(vaccine_rate=0.0)
731 print(round(cost, 2))

```

which gives

Python output

```

732 900000.0

```

Therefore without vaccinations, once the infection is endemic, the public health care system would expect to spend £900,000 a day.

With a vaccine rate of 85%:

Python input

```

733 cost = daily_cost(vaccine_rate=0.85)
734 print(round(cost, 2))

```

which gives

Python output

```

735 611903.36

```

So vaccinating 85% of newborns would cost the public health care system, once the infection is endemic £611,903.36 a day. That is a saving of around 32%.

5.4 SOLVING WITH R

In this book we will use the `deSolve` library to numerically solve the above epidemiology models.

We first define the system of differential equations described in Equations 5.1, 5.2 and 5.3. This is an R function where the arguments are the current time, the system state, and a list of other parameters, respectively.

R input

```

736 #' Defines the system of differential equations that describe
737 #' the epidemiology model.
738 #'
739 #' @param t a positive float
740 #' @param y a tuple of three integers
741 #' @param vaccine_rate a positive float <= 1
742 #' @param birth_rate a positive float <= 1
743 #'
744 #' @return a list containing dS, dI, and dR
745 derivatives <- function(t, y, parameters){
746   infection_rate <- 0.3
747   recovery_rate <- 0.02
748   death_rate <- 0.01
749   with(as.list(c(y, parameters)), {
750     N <- S + I + R
751     dSdt <- ( - ( (infection_rate * S * I) / N) # nolint
752               + ( (1 - vaccine_rate) * birth_rate * N)
753               - (death_rate * S))
754     dIdt <- ( ( (infection_rate * S * I) / N) # nolint
755               - (recovery_rate * I)
756               - (death_rate * I))
757     dRdt <- ( (recovery_rate * I) # nolint
758               - (death_rate * R)
759               + (vaccine_rate * birth_rate * N))
760     list(c(dSdt, dIdt, dRdt)) # nolint
761   })
762 }

```

Using this function returns the instantaneous rate of change for each of the three stocks, S , I and R . If we begin at time 0.0, with 4 susceptible individuals, 1 infected individual, 0 recovered individuals, a vaccine rate of 50% and a birth rate of 0.01, then:

R input

```
763 derivatives(t = 0,  
764             y = c(S = 4, I = 1, R = 0),  
765             parameters = c(vaccine_rate = 0.5,  
766                           birth_rate = 0.01)  
767 )
```

R output

```
768 [[1]]  
769 [1] -0.255  0.210  0.045
```

we would expect the number of susceptible individuals to reduce by around 0.255 per time unit, the number of infected individuals to increase by 0.21 per time unit, and the number of recovered individuals to increase by 0.045 per time unit. Now of course, after a tiny fraction of a time unit the stock levels will change, and thus the rates of change will change. So we will require something more sophisticated in order to determine the true behaviour of the system.

The following function observes the system's behaviour over some time period, using the `deSolve` library to numerically solve the system of differential equations:

R input

```

770 library(deSolve) # nolint
771
772 #' Numerically solve the system of differential equations
773 #'
774 #' @param t an array of increasing positive floats
775 #' @param y0 list of integers (default: c(S=2999, I=1, R=0))
776 #' @param birth_rate a positive float <= 1 (default: 0.01)
777 #' @param vaccine_rate a positive float <= 1 (default: 0.85)
778 #'
779 #' @return a matrix of times, S, I and R values
780 integrate_ode <- function(times,
781                             y0 = c(S = 2999, I = 1, R = 0),
782                             birth_rate = 0.01,
783                             vaccine_rate = 0.84){
784   params <- c(birth_rate = birth_rate,
785               vaccine_rate = vaccine_rate)
786   ode(y = y0,
787       times = times,
788       func = derivatives,
789       parms = params)
790 }

```

Now we can use this function to investigate the difference in behaviour between a vaccination rate of 0% and a vaccination rate of 85%. Let's observe the system for two years, that is 730 days, in time steps of 0.01 days.

Begin with a vaccine rate of 0%:

R input

```

791 times <- seq(0, 730, by = 0.01)
792 out <- integrate_ode(times, vaccine_rate = 0.0)

```

Now `out`, is a matrix with four columns, `time`, `S`, `I` and `R`, which are arrays of values of the time points, and the stock levels of `S`, `I` and `R` over the time respectively. We can plot these to visualise their behaviour. The following code gives the plot shown in Figure 5.4.

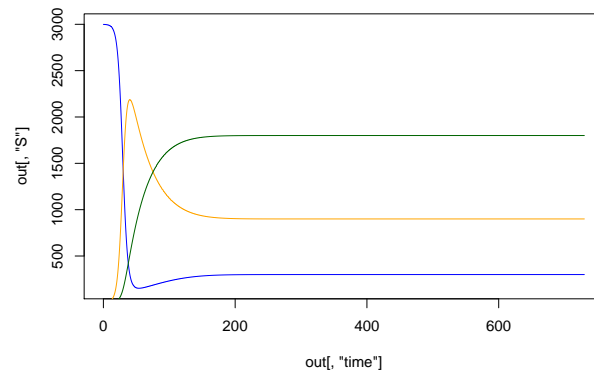


Figure 5.4 Output of code line 846-850

R input

```

793 pdf("plot_no_vaccine_R.pdf", width = 7, height = 5)
794 plot(out[, "time"], out[, "S"], type = "l", col = "blue")
795 lines(out[, "time"], out[, "I"], type = "l", col = "orange")
796 lines(out[, "time"], out[, "R"], type = "l", col = "darkgreen")
797 dev.off()

```

We observe that the number of infected individuals increases quickly, and in fact the rate of change increases as more individuals are infected. However this growth slows down as there are fewer susceptible individuals to infect. Due to the equal birth and death rates the overall population size remains constant; but we also see after some time period (around 300 time units) the levels of susceptible, infected, and recovered individuals becomes seemingly steady, and the disease becomes endemic. We can estimate once this steadiness occurs, around 10% of the population remain susceptible to the disease, 30% are infected, and 60% are recovered and immune.

Now with a vaccine rate of 85%:

R input

```

798 times <- seq(0, 730, by = 0.01)
799 out <- integrate_ode(times, vaccine_rate = 0.85)

```

And again we can plot these to visualise their behaviour. The following code gives the plot shown in Figure 5.5.

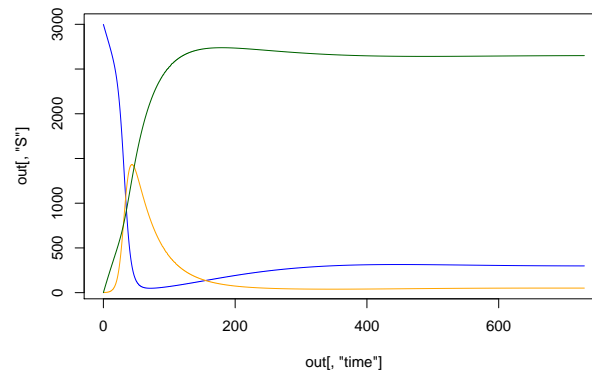


Figure 5.5 Output of code line 853-857

R input

```

800 pdf("plot_with_vaccine_R.pdf", width = 7, height = 5)
801 plot(out[, "time"], out[, "S"], type = "l", col = "blue")
802 lines(out[, "time"], out[, "I"], type = "l", col = "orange")
803 lines(out[, "time"], out[, "R"], type = "l", col = "darkgreen")
804 dev.off()

```

With vaccination the disease remains endemic, however now we estimate that once, steadiness occurs, around 10% of the population remain susceptible to the disease, 1.7% are infected, and 88.3% are immune or recovered and immune.

We've seen that vaccination lowers the percentage of the population living with the infection, which will lower the public healthcare system's medication costs. Let's now investigate if this saving is comparable to the cost of providing the vaccination to the newborns.

The following function calculates the total cost to the public healthcare system, that is the sum of the medication costs for those living with the infection and the vaccination costs:

R input

```

805 #' Calculates the daily cost to the public health
806 #' system after 2 years
807 #'
808 #' @param derivative_function: a function returning a
809 #'                               list of three floats
810 #' @param vaccine_rate: a positive float <= 1 (default: 0.85)
811 #'
812 #' @return the daily cost
813 daily_cost <- function(derivative_function = derivatives,
814                        vaccine_rate = 0.85){
815     max_time <- 730
816     time_step <- 0.01
817     birth_rate <- 0.01
818     vaccine_cost <- 220
819     medication_cost <- 10
820     times <- seq(0, max_time, by = time_step)
821     out <- integrate_ode(times, vaccine_rate = vaccine_rate)
822     N <- sum(tail(out[, c("S", "I", "R")], n = 1))
823     daily_vaccine_cost <- (N
824                           * birth_rate
825                           * vaccine_rate
826                           * vaccine_cost) / time_step
827     daily_medication_cost <- ( (tail(out[, "I"], n = 1)
828                              * medication_cost)) / time_step
829     daily_vaccine_cost + daily_medication_cost
830 }

```

Now let's compare the total daily cost with and without vaccination. Without vaccinations:

R input

```

831 cost <- daily_cost(vaccine_rate = 0.0)
832 print(cost)

```

which gives

R output

833

```
[1] 9e+05
```

Therefore without vaccinations, once the infection is endemic, the public health care system would expect to spend £900,000 a day.

With a vaccine rate of 85%:

R input

834

```
cost <- daily_cost(vaccine_rate = 0.85)  
print(cost)
```

835

which gives

R output

836

```
[1] 611903.4
```

So vaccinating 85% of newborns would cost the public health care system, once the infection is endemic £611,903.40 a day. That is a saving of around 32%.

5.5 RESEARCH



IV

Optimisation



Linear Programming

FINDING the best configuration of some system can be challenging, especially when there is a seemingly endless amount of possible solutions. Optimisation techniques are a way to mathematically derive solutions that maximise or minimise some objective function, subject to a number of feasibility constraints. When all components of the problem can be written in a linear way, then linear programming is one technique that can be used to find the solution.

6.1 PROBLEM

A university runs 26 modules over four subjects: Art, Biology, Chemistry, and Dutch. Each subject runs core modules and optional modules. Table 6.1 gives the module numbers for each of these.

The university is required to schedule examinations for each of these modules. The university would like the exams to be scheduled using the least amount of time slots possible. However not all modules can be scheduled at the same time as they share some students:

- All art modules share students,
- All biology modules share students,
- All chemistry modules share students,
- All Dutch modules share students,
- Art students have the option of taking the core Dutch modules, and so all art modules may share students with core Dutch modules,
- Biology students have the option of taking optional modules from chemistry, so all biology modules may share students with optional chemistry modules,
- Chemistry students have the option of taking optional modules from biology, so all chemistry modules may share students with optional biology modules,
- Biology students have the option of taking core art modules, and so all biology modules may share students with core art modules.

Art Core	Biology Core	Chemistry Core	Dutch Core
M00	M08	M16	M22
M01	M09	M17	M23
M02	M10	M18	
	M11	M19	
	M12		
Art Optional	Biology Optional	Chemistry Optional	Dutch Optional
M03	M13	M20	M24
M04	M14	M21	M25
M05	M15		
M06			
M07			

Table 6.1 List of modules on offer at the university.

What is the least number of exam time slots required to hold all 26 exams with no clashes?

6.2 THEORY

Linear programming is a method that solves an optimisation problem of n variables by defining all constraints as planes in n -dimensional space. These planes combine to create a convex region where all feasible solutions (those that satisfy the constraints) lie within the convex region, and all infeasible solutions (those that break at least one constraint) lie outside this convex region.

As we are interested in optimising some linear function, that is either minimising or maximising some function, the solution must lie at the very edge of the feasible convex region. That is we have improved so much that if we were to improve any further we would lie outside the feasible region - hence the optimum lies on the edge.

Linear programming employs algorithms such as the Simplex method to mathematically traverse the edges of the feasible convex region, and stops at the optimum. Therefore we only need to define our objective function and constraints in a linear fashion, and then apply appropriate algorithms.

Consider a 2-dimensional example: I am able to make £50 profit on each tonne of paint A I produce, and £60 profit on each tonne of paint B I produce. A tonne of paint A needs 4 tonnes of ingredient X and 5 tonnes of ingredient Y. A tonne of paint B needs 6 tonnes of ingredient X and 4 tonnes of ingredient Y. Only 24 tonnes of X and 20 tonnes of Y available per day. How much of paint A and paint B should I produce daily to maximise profit?

This is formulated as a linear objective function (total profit) to maximise, and two linear constraints (availability of ingredients X and Y). They are written as:

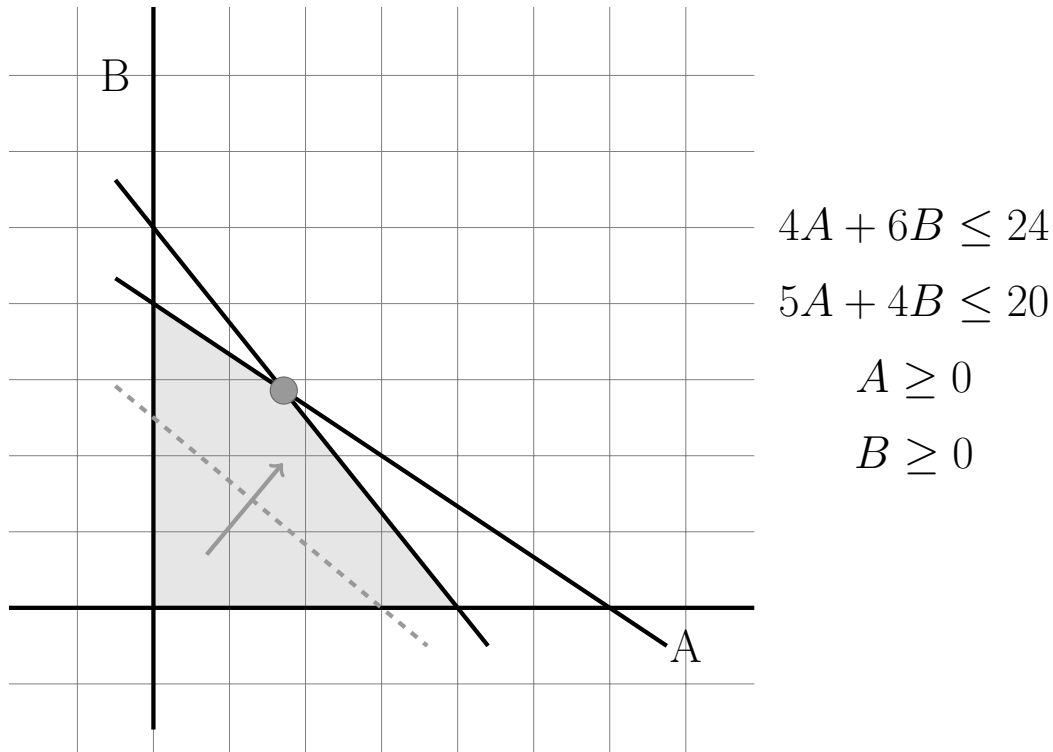


Figure 6.1 Visual representation of the paint linear program. The feasible convex region is shaded in grey; the objective function with arbitrary value is shown in a dashed line.

$$\text{Maximise: } 50A + 60B \quad (6.1)$$

Subject to:

$$4A + 6B \leq 24 \quad (6.2)$$

$$5A + 4B \leq 20 \quad (6.3)$$

Now we have a linear system in 2-dimensional space with coordinates A and B. These are called the decision variables, whose values we wish to find that optimises the objective given by expression 6.5. Inequalities 6.2 and 6.3 correspond to the amount of ingredient X and Y available per day. These, along with the additional constraints that we cannot produce a negative amount of paint ($A \geq 0$ and $B \geq 0$), form the convex feasible region shown in Figure 6.1.

Expression 6.5 corresponds to the total profit, which is the expression we are trying to maximise. As a line in the 2-dimensional space, this expression fixes its gradient, but its value determines the size of the y -intercept. Therefore optimising this function corresponds to pushing a line with that gradient to its furthest extreme




Figure 6.2 Visualisation of sets of modules with shared students.

within the feasible region, demonstrated in Figure ???. Therefore for this problem the optimum occurs in a particular vertex of the feasible region, at $A = \frac{12}{7}$ and $B = \frac{20}{7}$.

This works well as A and B can take any real value in the feasible region. It is common however to formulate Integer Linear Programmes where the decision variables are restricted to integers. There are a number of methods that can help us adapt a real solution to an integer solution. These include cutting planes, which introduce new constraints around the real solution to force an integer value; and branch and bound methods, where we iteratively convert decision variables to their closest two integers and remove any infeasible solutions.

Both Python and R have libraries that carry out the linear and integer programming algorithms for us. When solving these kinds of problems, formulating them as linear systems is the most important challenge.

Consider again the exam scheduling problem from Section 6.1. There are eight distinct sets of modules in which exams cannot be scheduled simultaneously: A_c , A_o representing core and optional art modules respectively; B_c , B_o representing core and optional biology modules respectively; C_c , C_o representing core and optional chemistry modules respectively; and D_c , D_o representing core and optional Dutch modules respectively.

Additionally there are further clashes between these sets, this can be visualised as a graph, shown in Figure 6.2 where edges between sets represent shared students. This shows there are five cliques (shown in red, blue, green, orange and yellow), that is five larger sets of modules that may share students: $C_o \cup C_c \cup B_o$, $C_o \cup B_o \cup B_c$, $B_c \cup B_o \cup A_c$, $A_c \cup A_o \cup D_c$, and $D_c \cup D_o$. These sets of modules will form the basis for our constraints, Inequalities 6.8 to 6.12.

Define M as the set of all modules to be scheduled, that is $M = A_c \cup A_o \cup B_c \cup B_o \cup C_c \cup C_o \cup D_c \cup D_o$. Define also T as the set of possible time slots. At worst each exam is scheduled for a different day, thus $|T| = |M| = 26$ in this case. Let $\{X_{mt} \text{ for } m \in M \text{ and } t \in T\}$ be a set of binary decision variables, that is X_{mt} is 1 if module m is scheduled for time t , and 0 otherwise. Let's also define $\{Y_t \text{ for } t \in T\}$ as a set of auxiliary binary decision variables, where Y_t is 1 if time slot t is being used. This is enforced by Inequality 6.6.

Finally we have one final constraint, Inequality 6.7, which ensures all modules are scheduled once and once only. Thus altogether our integer program becomes:

$$\text{Minimise: } \sum_{t \in T} Y_j \quad (6.4)$$

Subject to:

$$\frac{1}{|M|} \sum_{m \in M} X_{mt} \leq Y_j \text{ for all } j \in T \quad (6.5)$$

$$\sum_{t \in T} X_{mt} = 1 \text{ for all } m \in M \quad (6.6)$$

$$\sum_{m \in C_o \cup C_c \cup B_o} X_{mt} \leq 1 \text{ for all } t \in T \quad (6.7)$$

$$\sum_{m \in C_o \cup B_o \cup B_c} X_{mt} \leq 1 \text{ for all } t \in T \quad (6.8)$$

$$\sum_{m \in B_c \cup B_o \cup A_c} X_{mt} \leq 1 \text{ for all } t \in T \quad (6.9)$$

$$\sum_{m \in A_c \cup A_o \cup D_c} X_{mt} \leq 1 \text{ for all } t \in T \quad (6.10)$$

$$\sum_{m \in D_c \cup D_o} X_{mt} \leq 1 \text{ for all } t \in T \quad (6.11)$$

$$(6.12)$$

6.3 SOLVING WITH PYTHON

In this book we will use the Python library PuLP to formulate and solve the integer program. First let's define all the sets we will use to formulate the problem.

Python input

```

837 Ac = [0, 1, 2]
838 Ao = [3, 4, 5, 6, 7]
839 Bc = [8, 9, 10, 11, 12]
840 Bo = [13, 14, 15]
841 Cc = [16, 17, 18, 19]
842 Co = [20, 21]
843 Dc = [22, 23]
844 Do = [24, 25]
845 modules = Ac + Ao + Bc + Bo + Cc + Co + Dc + Do
846 time_slots = range(26)

```

Now let's begin by defining an empty problem:

Python input

```
847 import pulp
848 prob = pulp.LpProblem("ExamScheduling", pulp.LpMinimize)
```

We also need to define our sets of binary decision variables:

Python input

```
849 x = pulp.LpVariable.dicts("X", (modules, time_slots), cat=pulp.LpBinary)
850 y = pulp.LpVariable.dicts('Y', time_slots, cat=pulp.LpBinary)
```

Now y is a dictionary of binary decision variables, with keys as elements of the list `time_slots`. Let's look at Y_3 corresponding to the third day:

Python input

```
851 print(y[3])
```

Python output

```
852 Y_3
```

While x is a dictionary of dictionaries of binary decision variables, with keys as elements of the lists `modules` and `time_slots`. Let's look at $X_{2,5}$, the variable corresponding to module 2 being scheduled on day 5:

Python input

```
853 print(x[2][5])
```

Python output

```
854 X_2_5
```

Now we have an empty problem, all relevant sets, and all decision variables defined, we can go ahead and add the objective function and constraints to the problem.

For the objective function:

Python input

```
855 objective_function = sum([y[day] for day in time_slots])
856 prob += objective_function
```

Now the constraints:

Python input

```
857 for day in time_slots:
858     prob += (1/len(modules)) * sum([x[module][day] for module in modules]) <= y[day]
859     prob += sum([x[module][day] for module in Ac+Ao+Dc]) <= 1
860     prob += sum([x[module][day] for module in Bc+Bo+Co]) <= 1
861     prob += sum([x[module][day] for module in Bc+Bo+Ac]) <= 1
862     prob += sum([x[module][day] for module in Cc+Co+Bo]) <= 1
863     prob += sum([x[module][day] for module in Dc+Do]) <= 1
864
865 for module in modules:
866     prob += sum([x[module][day] for day in time_slots]) == 1
```

At this stage we could print the `prob` object, which would explicitly give all constraints written out fully. This can be used to error check if the need arises.

Now we can go ahead and solve the problem:

Python input

```
867 prob.solve()
```

This returns a status code: 1 for 'optimal', 0 for 'not solved', -1 for 'infeasible', -2 for 'unbounded', or -3 for 'undefined'. A problem is successfully solved if this method returns 1:

Python output

```
868 1
```

This method has also assigned values to our decision variables. These can be inspected, let's check if module 2 was scheduled for day 5:

Python input

```
869 print(x[2][5])
```

Python output

```
870 0.0
```

This was assigned the value 0, and so module 2 was not scheduled for that day. Let's check if module 2 was scheduled for day 14:

Python input

```
871 print(x[2][14])
```

Python output

```
872 1.0
```

This was assigned a value of 1, and so module 2 was scheduled for that day.

We can iterate through all decision variables and make a print solutions in order to read off the schedule easier:

Python input

```
873 for day in time_slots:
874     if y[day].value() == 1:
875         schedule = f"Day {day}: "
876         for module in modules:
877             if x[module][day].value() == 1:
878                 schedule += f"{module}, "
879         print(schedule)
```

giving:

Python output

```

880 Day 0: 7, 9, 25,
881 Day 1: 6, 12, 17, 24,
882 Day 3: 3, 14,
883 Day 4: 10, 19,
884 Day 10: 4, 8, 16,
885 Day 11: 13, 22,
886 Day 12: 11, 23,
887 Day 13: 5, 15,
888 Day 14: 2, 20,
889 Day 17: 0, 18,
890 Day 18: 1, 21,

```

Now the order of the days do not matter here, but we can see that 11 days are required in order to schedule all exams with no clashes. Most days (8) have two exams scheduled, 2 days have 3 exams scheduled, and it was possible on 1 day to have 4 exams scheduled simultaneously.

6.4 SOLVING WITH R

In R we will use the R package `ROI`, the R Optimization Infrastructure. This is a library of code that acts as a front end to a number of other solvers that need to be installed externally, allowing a range of optimisation problems to be solved with a number of different solvers, using similar problem structures and syntax. The solver that we will use here is called the CBC MILP Solver, which needs to be installed as well as `rcbc` package.

In Section 6.2 the exam scheduling problem was formulated with two arrays of decision variables, a one-dimensional array Y and a two dimensional array X . In `ROI` we need our decision variables to be a one-dimensional array, so we ‘flatten’ these by first ordering by X then Y , within that ordering by time slot, then within that ordering by module number. Call these variables Z_i for $i \in 1, 2, \dots, 702$. For example Z_{29} would correspond to $X_{3,2}$, the decision variable representing whether module number 3 is scheduled on day 2; Z_{700} would correspond to Y_{24} , the decision variable representing whether there’s an exam scheduled for day 24.

Once this array is defined, then the objective function is an ordered array of the coefficients of each of these variables:

R input

```

891  #' Writes the row of coefficients for the objective function
892  #'
893  #' @param n_modules: the number of modules to schedule
894  #' @param n_days: the maximum number of days to schedule
895  #'
896  #' @return the objective function row to minimise
897  write_objective <- function(n_modules, n_days){
898    all_days <- rep(0, n_modules * n_days)
899    Ys <- rep(1, n_days)
900    append(all_days, Ys)
901  }

```

For 3 modules and 3 days:

R input

```

902  write_objective(3, 3)

```

Which gives the following array, corresponding the the coefficients of the array Z for Equation 6.5.

R output

```

903  [1] 0 0 0 0 0 0 0 0 0 1 1 1

```

Now the constraints have three components: the coefficients matrix, the right hand side array, and the array of directions. The coefficients matrix is a matrix of the coefficients of the left hand side of inequalities, where again columns represent the coefficients of the array Z , and rows represent each separate constraint. The following function is used to write one row of that coefficients matrix, for a given day, for a given set of clashes, corresponding to Inequalities 6.8 to 6.12:

R input

```

904  #' Writes the constraint row dealing with clashes
905  #'
906  #' @param clashes: a vector of module indices that all cannot
907  #'                  be scheduled at the same time
908  #' @param day: an integer representing the day
909  #'
910  #' @return the constraint row corresponding to that set of
911  #'         clashes on that day
912  write_X_clashes <- function(clashes, day, n_days, n_modules){
913    today <- rep(0, n_modules)
914    today[clashes] = 1
915    before_today <- rep(0, n_modules * (day - 1))
916    after_today <- rep(0, n_modules * (n_days - day))
917    all_days <- c(before_today, today, after_today)
918    full_coeffs <- c(all_days, rep(0, n_days))
919    full_coeffs
920  }

```

where an array containing the module numbers of a set of modules that may all share students. The following function is used to write one row of the coefficients matrix, for each module, ensuring that each module is scheduled on one day and one day only, corresponding to Equation 6.7:

R input

```

921  #' Writes the constraint row to ensure that every module is
922  #' scheduled once and only one
923  #'
924  #' @param module: an integer representing the module
925  #'
926  #' @return the constraint row corresponding to scheduling a
927  #'         module on only one day
928  write_X_requirements <- function(module, n_days, n_modules){
929    today <- rep(0, n_modules)
930    today[module] = 1
931    all_days <- rep(today, n_days)
932    full_coeffs <- c(all_days, rep(0, n_days))
933    full_coeffs
934  }

```

The following function is used to write one row of the coefficients matrix corresponding to the auxiliary constraints of Inequalities ??:

R input

```

935  #' Writes the constraint row representing the Y variable,
936  #' whether at least one exam is scheduled on that day
937  #'
938  #' @param day: an integer representing the day
939  #'
940  #' @return the constraint row corresponding to creating Y
941  write_Y_constraints <- function(day, n_days, n_modules){
942    today <- rep(1, n_modules)
943    before_today <- rep(0, n_modules * (day - 1))
944    after_today <- rep(0, n_modules * (n_days - day))
945    all_days <- c(before_today, today, after_today)
946    all_Ys <- rep(0, n_days)
947    all_Ys[day] = -n_modules
948    full_coeffs <- append(all_days, all_Ys)
949    full_coeffs
950  }

```

Finally the following function uses them all to assemble a coefficients matrix. It loops through the parameters for each constraint row required, uses the appropriate function to create the row of the coefficients matrix, sets the appropriate inequality direction (\leq , $=$, \geq), and the value of the right hand side. It returns all three components:

R input

```

951 #' Writes all the constraints as a matrix, column of
952 #' inequalities, and right hand side column.
953 #'
954 #' @param list_clashes: a list of vectors with sets of modules
955 #' that cannot be scheduled at the same time
956 #'
957 #' @return f.con the LHS of the constraints as a matrix
958 #' @return f.dir a vector of directions of the inequalities
959 #' @return f.rhs a vector of the values of the RHS of the inequalities
960 write_constraints <- function(list_clashes, n_days, n_modules){
961   all_rows <- c()
962   all_dirs <- c()
963   all_rhss <- c()
964   n_rows <- 0
965
966   for (clash in list_clashes){
967     for (day in 1:n_days){
968       clashes <- write_X_clashes(clash, day, n_days, n_modules)
969       all_rows <- append(all_rows, clashes)
970       all_dirs <- append(all_dirs, "<=")
971       all_rhss <- append(all_rhss, 1)
972       n_rows <- n_rows + 1
973     }
974   }
975
976   for (module in 1:n_modules){
977     reqs <- write_X_requirements(module, n_days, n_modules)
978     all_rows <- append(all_rows, reqs)
979     all_dirs <- append(all_dirs, "==")
980     all_rhss <- append(all_rhss, 1)
981     n_rows <- n_rows + 1
982   }
983
984   for (day in 1:n_days){
985     Yconstraints <- write_Y_constraints(day, n_days, n_modules)
986     all_rows <- append(all_rows, Yconstraints)
987     all_dirs <- append(all_dirs, "<=")
988     all_rhss <- append(all_rhss, 0)
989     n_rows <- n_rows + 1
990   }
991
992   f.con <- matrix(all_rows, nrow = n_rows, byrow = TRUE)
993   f.dir <- all_dirs
994   f.rhs <- all_rhss
995   list(f.con, f.dir, f.rhs)
996 }

```

For demonstration, if we had two modules and two possible days, with the single constraint that both modules cannot be scheduled at the same time, then:

R input

```

997 write_constraints(list_clashes = list(c(1, 2)),
998                   n_days = 2,
999                   n_modules = 2)

```

This would give three components, a coefficient matrix of the left hand side of the constraints (rows 1 and 2 corresponding to the clash on days 1 and 2, row 3 ensuring module 1 is scheduled on one day only, row 4 ensuring module 2 is scheduled on one day only, and rows 5 and 5 defining the Y decision variables), an array of direction of the constraint inequalities, and an array of the right hand side values of the constraints.

R output

```

1000 [[1]]
1001      [,1] [,2] [,3] [,4] [,5] [,6]
1002 [1,]    1    1    0    0    0    0
1003 [2,]    0    0    1    1    0    0
1004 [3,]    1    0    1    0    0    0
1005 [4,]    0    1    0    1    0    0
1006 [5,]    1    1    0    0   -2    0
1007 [6,]    0    0    1    1    0   -2
1008
1009 [[2]]
1010 [1] "<=" "<=" "==" "==" "<=" "<="
1011
1012 [[3]]
1013 [1] 1 1 1 1 0 0

```

Now we are ready to use these to solve the exam scheduling problem. First we define some parameters, including the sets of modules that all share students, that is the list of clashes:

R input

```

1014 n_modules = 26
1015 n_days = 26
1016
1017 Ac <- c(1, 2, 3)
1018 Ao <- c(4, 5, 6, 7, 8)
1019 Bc <- c(9, 10, 11, 12, 13)
1020 Bo <- c(14, 15, 16)
1021 Cc <- c(17, 18, 19, 20)
1022 Co <- c(21, 22)
1023 Dc <- c(23, 24)
1024 Do <- c(25, 26)
1025
1026 list_clashes <- list(
1027   c(Ac, Ao, Dc),
1028   c(Bc, Bo, Co),
1029   c(Bc, Bo, Ac),
1030   c(Cc, Co, Bo),
1031   c(Dc, Do)
1032 )

```

Then we can use the functions defined above to create the objective function and the three elements of the constraints:

R input

```

1033 constraints <- write_constraints(list_clashes = list_clashes,
1034                                n_days = n_days,
1035                                n_modules = n_modules)
1036 f.con <- constraints[[1]]
1037 f.dir <- constraints[[2]]
1038 f.rhs <- constraints[[3]]
1039 f.obj <- write_objective(n_modules = n_modules, n_days = n_days)

```

Finally, once these objects are in place, we can use the to construct an optimisation problem object:

R input

```

1040 library(ROI)
1041
1042 milp <- OP(objective = L_objective(f.obj),
1043           constraints = L_constraint(L = f.con,
1044                                   dir = f.dir,
1045                                   rhs = f.rhs),
1046           types = rep("B", length(f.obj)),
1047           maximum = FALSE)

```

This creates an OP object from our objective row `f.obj`, and our constraints which are made up from the three components `f.con`, `f.dir` and `f.rhs`. When creating this object we also denote the `types` as binary variables (an array of "B" for each decision variable), and we want to minimise our objective function so we set `maximum = FALSE`.

Now to solve:

R input

```

1048 sol <- ROI_solve(milp)

```

The solver will not output information about the solve process and runtime. We can now print the solution:

R input

```

1049 print(sol$solution)

```

R output

```

1050 [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1051 [30] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1052 [59] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1053 [88] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1054 [117] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
1055 [146] 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1056 [175] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1
1057 [204] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1058 [233] 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 0 0 0
1059 [262] 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0
1060 [291] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1061 [320] 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0
1062 [349] 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1063 [378] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1064 [407] 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0
1065 [436] 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0
1066 [465] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1067 [494] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1068 [523] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1069 [552] 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0
1070 [581] 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0
1071 [610] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
1072 [639] 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1073 [668] 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 1 0 1 1 0 0 1 1
1074 [697] 0 1 1 0 1 0

```

This gives the values of each of the Z decision variables. We know the structure of this, that is the first 26 variables are the modules scheduled for day 1, and so on. The following code prints a readable schedule:

R input

```

1075 for (day in 1:n_days){
1076   if (sol$solution[(n_days*n_modules) + day] == 1){
1077     schedule <- paste("Day", day, ":")
1078     for (module in 1:n_modules){
1079       var <- ((day - 1) * n_modules) + module
1080       if (sol$solution[var] == 1){
1081         schedule <- paste(schedule, module)
1082       }
1083     }
1084     print(schedule)
1085   }
1086 }

```

R output

```

1087 [1] "Day 6 : 2 22"
1088 [1] "Day 8 : 8 9 19"
1089 [1] "Day 10 : 5 13 20"
1090 [1] "Day 11 : 3 21 26"
1091 [1] "Day 13 : 16"
1092 [1] "Day 14 : 6 15"
1093 [1] "Day 17 : 4 14 25"
1094 [1] "Day 18 : 7 11"
1095 [1] "Day 22 : 12 17 24"
1096 [1] "Day 23 : 10 23"
1097 [1] "Day 25 : 1 18"

```

Giving that 11 days are the minimum required to schedule the 26 exams without clashes.

6.5 RESEARCH

Bibliography

